

Homework 1: Introduction to Numpy

Collaborators

Please list anyone you discussed or collaborated on this assignment with below.

LIST COLLABORATORS HERE

Python setup

```
# Run me first!
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
sns.set_style('darkgrid')
```

Part 1: Numpy basics

As discussed in class, a square matrix A defines a linear mapping: $\mathbb{R}^n \rightarrow \mathbb{R}^n$. Given a vector \mathbf{x} , we can find the corresponding output of this mapping \mathbf{b} using matrix-vector multiplication: $\mathbf{b} = A\mathbf{x}$. We can write an example matrix-multiplication using matrix notation as:

$$\begin{bmatrix} 4 & -3 & 2 \\ 6 & 5 & 1 \\ -4 & -1 & 2 \end{bmatrix} \cdot \begin{bmatrix} 1 \\ -2 \\ 1 \end{bmatrix} = \begin{bmatrix} ? \\ ? \\ ? \end{bmatrix}$$

Q1

Perform this matrix-vector multiplication by hand and write the answer in the cell below.

WRITE ANSWER HERE

Q2

In the code cell below, create the matrix A and the vector \mathbf{x} shown above, using Numpy. Then use the `np.dot` function to find the output of the mapping $\mathbf{b} = A\mathbf{x}$. Verify that the answer matches what you derived above.

```
# Fill answers here
A =
x =
b =

print(b)
```

Often we will have access to the transformed vector \mathbf{b} and need to find the original vector \mathbf{x} . To do this we need to *solve* the system of linear equations $A\mathbf{x} = \mathbf{b}$ for \mathbf{x} .

$$\begin{bmatrix} 4 & -3 & 2 \\ 6 & 5 & 1 \\ -4 & -1 & 2 \end{bmatrix} \cdot \begin{bmatrix} ? \\ ? \\ ? \end{bmatrix} = \begin{bmatrix} 2 \\ -1 \\ 3 \end{bmatrix}$$

Q3

Find the missing \mathbf{x} in the equation above using the `np.linalg.solve` function and verify that $A\mathbf{x} = \mathbf{b}$.

```
# Fill answer here (A is the same matrix from above)
b =
x =
```

In linear algebra you may have learned how to solve a system of linear equations using Gaussian elimination. Here we will implement an alternative approach known as *Richardson iteration*. In this method we start with an initial guess for the solution: $\mathbf{x}^{(0)}$, then we will iteratively update this guess until the solution is reached. Given a matrix A , a target \mathbf{b} and a current guess $\mathbf{x}^{(k)}$, we can compute the Richardson update as:

$$\mathbf{x}^{(k+1)} \leftarrow \mathbf{x}^{(k)} + \omega(\mathbf{b} - A\mathbf{x}^{(k)})$$

Here ω is a constant that we can choose to adjust the algorithm. We will set $\omega = 0.1$.

Q4

Fill in the Richardson iteration function below and apply it to the system of linear equations from above using 100 updates. Verify that it gives a similar answer to `np.linalg.solve`.

```
# Fill in function below
def richardson_iter(x_guess, A, b, omega=0.1):

    return new_x_guess

x_guess = np.zeros(3)
for i in range(100):
    x_guess = richardson_iter(x_guess, A, b)

print(x_guess, x)
```

Recall that the length of a vector is given by its *two-norm*, which is defined as:

$$\|\mathbf{x}\|_2 = \sqrt{\sum_{i=1}^n x_i^2}.$$

Correspondingly, the (Euclidian) distance between two points $\mathbf{a}, \mathbf{b} \in \mathbb{R}^n$ can be written as $\|\mathbf{a} - \mathbf{b}\|_2$. As a convenient measure of error for our Richardson iteration algorithm, we will use the *squared Euclidean distance*. For a guess $\mathbf{x}^{(k)}$ we will compute the error $e^{(k)}$ as:

$$e^{(k)} = \|A\mathbf{x}^{(k)} - \mathbf{b}\|_2^2$$

In expanded form, this would be written as:

$$e^{(k)} = \sum_{i=1}^n \left(\sum_{j=1}^n A_{ij} x_j^{(k)} - b_i \right)^2$$

Q5

Write a function to compute the error of a given guess. Then run Richardson iteration again for 100 steps, computing the error at each step. Finally create a plot of the error for each step (error vs. step). [Plot reference](#)

Hint: recall that basic operations in numpy (addition, subtraction, powers) are performed element-wise.

```
# Fill in function below
def error(x_guess, A, b):

    return err

# Add code to plot the error over time

x_guess = np.zeros(3)
for step in range(100):

    x_guess = richardson_iter(x_guess, A, b)
```

Q6

Derive the partial derivative of the error with respect to a single entry of $\mathbf{x}^{(k)}$ (without loss of generality, we will say $x_1^{(k)}$). Work in the *expanded form* as in the equation above, writing your answer in the markdown cell below.

Hint: You may find it helpful to refer to the latex equation cheatsheet on the course website. You may show intermediate steps here or as handwritten work as a separate file in the repository. The final answer should be filled in here.

EDIT THIS CELL WITH YOUR ANSWER

$$\frac{\partial e^{(k)}}{\partial x_1^{(k)}} =$$

YOU MAY ADD WORK HERE

In practice, we will likely want to compute the derivative with respect to *all* entries of \mathbf{x} :

$$\frac{\partial e^{(k)}}{\partial \mathbf{x}^{(k)}} = \begin{bmatrix} \frac{\partial e^{(k)}}{\partial x_1^{(k)}} & \dots & \frac{\partial e^{(k)}}{\partial x_n^{(k)}} \end{bmatrix}$$

Q7

Using the formula you just derived, write the formula for the vector of all partial derivatives in the compact matrix/vector notation (e.g. $A\mathbf{x} = \mathbf{b}$).

EDIT THIS CELL WITH YOUR ANSWER

$$\frac{\partial e^{(k)}}{\partial \mathbf{x}^{(k)}} =$$

Q8

Do you notice any relationship between this result and the Richardson iteration algorithm above? (1-2 sentences) *We will discuss this more in class!*

WRITE YOUR ANSWER HERE

Part 2: Working with batches of vectors

Recall that a vector can also be seen as either an $n \times 1$ matrix (column vector) or a $1 \times n$ matrix (row vector).

$$\begin{aligned} \text{Column vector: } \mathbf{x} &= \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}, \\ \text{Row vector: } \mathbf{x}^T &= \begin{bmatrix} x_1 & x_2 & x_3 \end{bmatrix} \end{aligned}$$

Note that we use the same notation for both as they refer to the same concept (a vector). The difference becomes relevant when we consider matrix-vector multiplication. We can write

matrix-vector multiplication in two ways:

$$\text{Matrix-vector: } A\mathbf{x} = \mathbf{b}, \quad \text{Vector-matrix: } \mathbf{x}^T A^T = \mathbf{b}^T$$

In *matrix-vector multiplication* we treat \mathbf{x} as a column vector ($n \times 1$ matrix), while in *vector-matrix multiplication* we treat it as a row vector ($1 \times n$ matrix). Transposing A for left multiplication ensures that the two forms give the same answer.

Q9

Using the previously defined \mathbf{x} , create an explicit column vector and row vector. Then using the previously defined A , verify that the matrix-vector and vector-matrix multiplications shown above do produce the same resultant vector \mathbf{b} .

Hint: Recall that `np.dot` is also used for matrix-matrix multiplication. The values of \mathbf{x} and A are:

$$\mathbf{x} = \begin{bmatrix} 1 \\ -2 \\ 1 \end{bmatrix}, \quad A = \begin{bmatrix} 4 & -3 & 2 \\ 6 & 5 & 1 \\ -4 & -1 & 2 \end{bmatrix}$$

```
# Fill in code here
x_col =
x_row =
```

The distinction between row and column vectors also affects the behavior of other operations on `np.array` objects, particularly through the concept of [broadcasting](#).

Q10

Consider a 3×3 matrix of all ones as defined in code below, along with the 1-d vector \mathbf{x} .

```
ones = np.ones((3, 3))
x = np.array([1, -2, 1])
ones
```

```
array([[1., 1., 1.],
       [1., 1., 1.],
       [1., 1., 1.]])
```

Complete the line of code below such that the result of the operation is:

$$\begin{bmatrix} 1 & -2 & 1 \\ 1 & -2 & 1 \\ 1 & -2 & 1 \end{bmatrix}$$

Hint: You should replace `SHAPE` with an appropriate shape for broadcasting and `OPERATION` with an appropriate operation (e.g. `+`, `-`, ``, `/`)*

```
# Fill in code here
ones OPERATION x.reshape( SHAPE )
```

Throughout this course we will typically use row vectors and vector-matrix multiplication, as this is more conventional in neural-network literature. The concept of row and column vectors becomes handy when transforming *collections* of vectors.

Recall that a matrix can be seen as a collection of vectors. In numpy we can create a matrix from a list of (1-dimensional) vectors using the `np.stack` function. This function assumes that the vectors are row vectors creating the matrix as follows:

$$\begin{bmatrix} 3 & 1 & -2 \\ -2 & -1 & 5 \end{bmatrix} \quad \text{vs} \quad \begin{bmatrix} 4 & 5 & 3 \\ 3 & 1 & -2 \end{bmatrix}$$

$\xrightarrow{\text{np.stack}}$

$$\begin{bmatrix} 3 & 1 & -2 & 4 & 5 & 3 \\ -2 & -1 & 5 & 3 & 1 & -2 \end{bmatrix}$$

We will call this matrix X to denote that it is a collection of vectors, rather than a single vector (\mathbf{x}).

Q11

Create this matrix in numpy using the `np.stack` function.

```
# Fill code here
```

```
X =  
print(X)
```

When taken together as a matrix in this way, we can apply the linear mapping A to all vectors using matrix-matrix multiplication:

$$B = XA^T$$

Let's put this into practice with a visual example.

Q12

Create a 20×3 matrix, `circle`, in numpy of the following form

$$\begin{bmatrix} \sin(\theta_1) & \cos(\theta_1) & 1 \\ \sin(\theta_2) & \cos(\theta_2) & 1 \\ \vdots & \vdots & \vdots \\ \sin(\theta_{20}) & \cos(\theta_{20}) & 1 \end{bmatrix}$$

Where $\theta_1 \dots \theta_{20}$ are evenly spaced between 0 and 2π .

```
theta = np.linspace(0, 2 * np.pi, 20) # Generates 20 evenly-spaced numbers  
# Fill in your code here  
circle =
```

The code we just wrote creates a matrix corresponding to a collection of 20 row vectors of length 3. Each vector represents a point on the unit circle where the first entry is the x-coordinate, the second entry is the y-coordinate and the third entry is always 1:

$$\begin{bmatrix} x & y & 1 \end{bmatrix}$$

Q13

Plot the set of 20 points in `circle` using the `plt.plot` function. *Use only the x and y coordinates, ignoring the column of 1s.*

```
plt.figure(figsize=(4, 4))  
plt.xlim(-3, 3)  
plt.ylim(-3, 3)  
# Fill your code here
```

Q14

Transform all the vectors in `circle` with the matrix A using a single call to `np.dot`. Then plot the original set of points in black and the transformed points in red using the `plt.plot` function.

You might also consider why we added the extra column of 1s! We will discuss the answer to that in class. A is the same matrix from q1.

```
# Fill your code here  
transformed_circle =  
  
plt.figure(figsize=(4, 4))  
plt.xlim(-3, 3)  
plt.ylim(-3, 3)  
# Fill your code here
```

Q15

Finally, shift all the vectors in `transformed_circle` by the vector **b** defined below, that is, add **b** to every *row* of the output matrix. Again plot the original set of points in black and the transformed points in red using the `plt.plot` function.

*Hint: the solution to this question should **not** involve any loops, instead use broadcasting.*

```
# Fill your code here
b = np.array([-0.5, 1.2, 0])
transformed_circle =
```

```
plt.figure(figsize=(4, 4))
plt.xlim(-3, 3)
plt.ylim(-3, 3)
```

```
# Fill your code here
```

Part 3: Loading and visualizing data

For most of this class we will be working with real-world data. A very well-known dataset in statistics is the [Iris flower dataset](#) collected by Edgar Anderson in 1929. The dataset consists of measurements of iris flowers. Each flower has 4 collected measurements: sepal length, sepal width, petal length, petal width, as well as a classification into one of 3 species: *Iris setosa*, *Iris versicolor* and *Iris virginica*. We will return to this dataset in the next homework.

We can load this dataset as Numpy objects using the Scikit-Learn library. Below we've extracted 4 relevant arrays: - **features**: a matrix which has one row per observed flower and one column per measurement. - **targets**: An array that specifies the species of each flower as a number 0-2. - **feature_names**: a list of strings with the name of each measurement. - **target_names**: a list of strings with the name of each species.

In this homework, we will only visualize this dataset, which is typically a good first step in working with a new type of data. We'll start by just looking at 2 measurements *sepal length* and *petal length*, along with the species.

Q16

Based on the Iris dataset loaded below, how many flowers did Edgar Anderson measure?

In other words, how many observations are in this dataset?

```
import sklearn.datasets as datasets
dataset = datasets.load_iris()
features = dataset['data']
targets = dataset['target']
feature_names = dataset['feature_names']
target_names = dataset['target_names']
```

```
# Write any code to determine the number of observations here
```

Fill in the code to create a scatterplot for the Iris dataset below. Plot *sepal length* on the **x-axis** and *petal length* on the **y-axis**. Set the **color** to correspond to the *species*.

```
# Fill your code here

plt.figure(figsize=(4, 4))
```