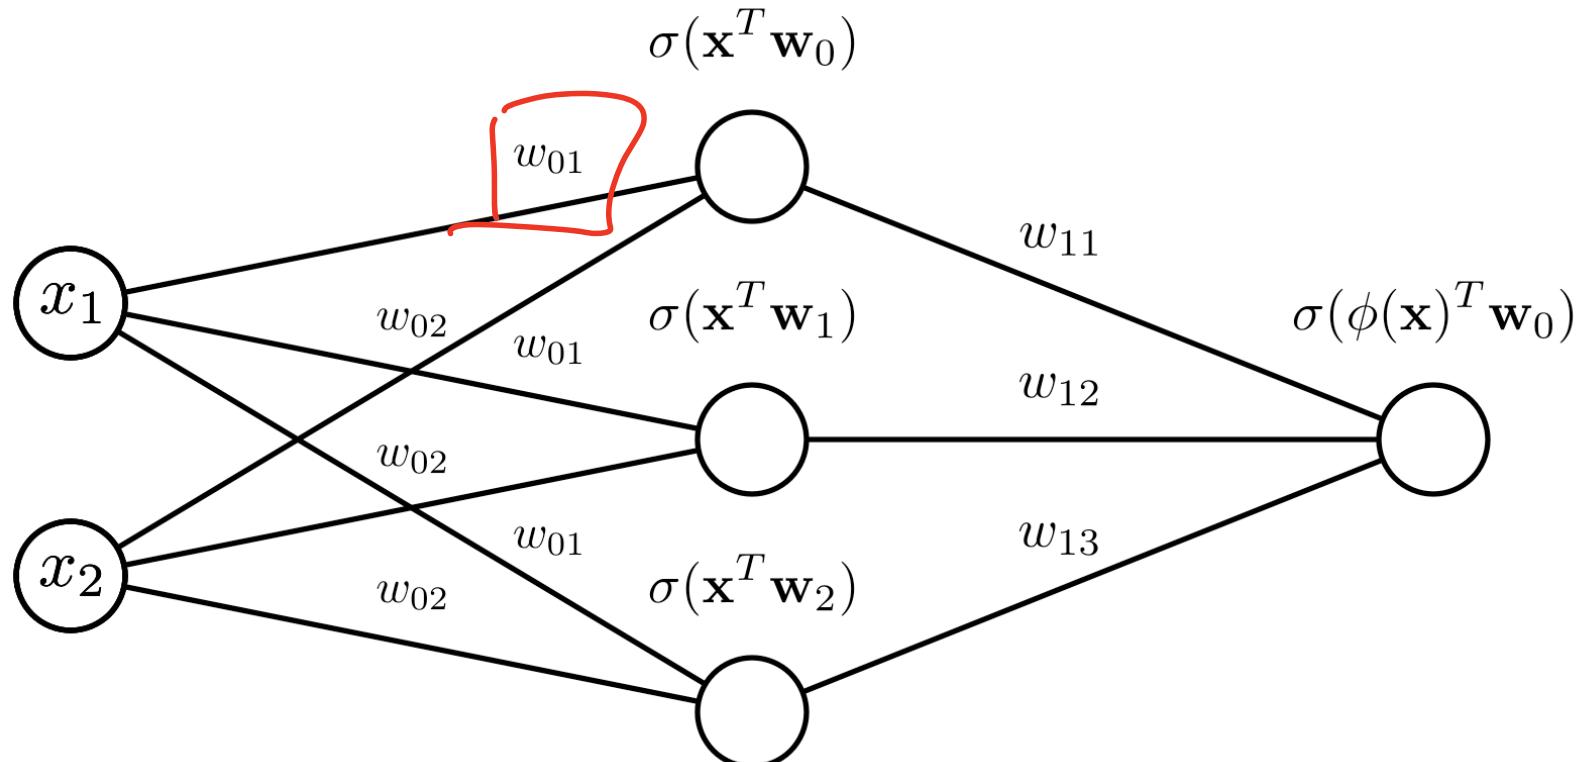
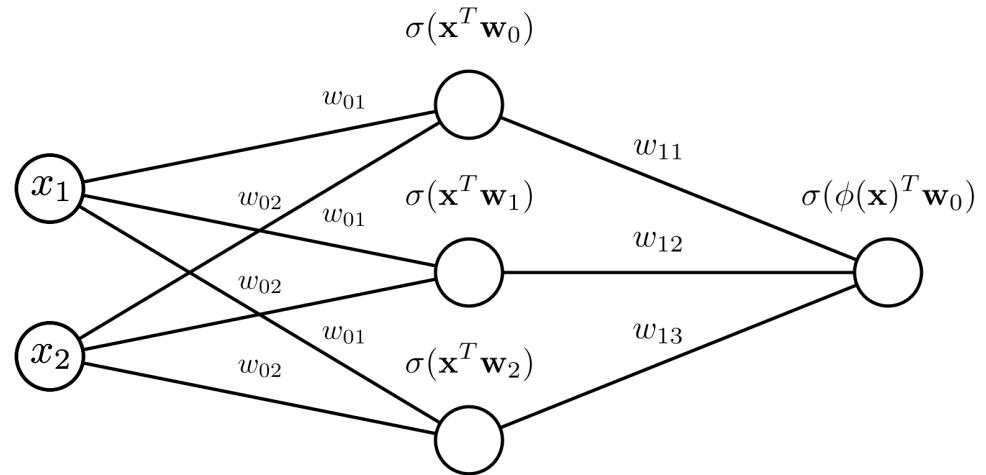


Learning Neural Networks

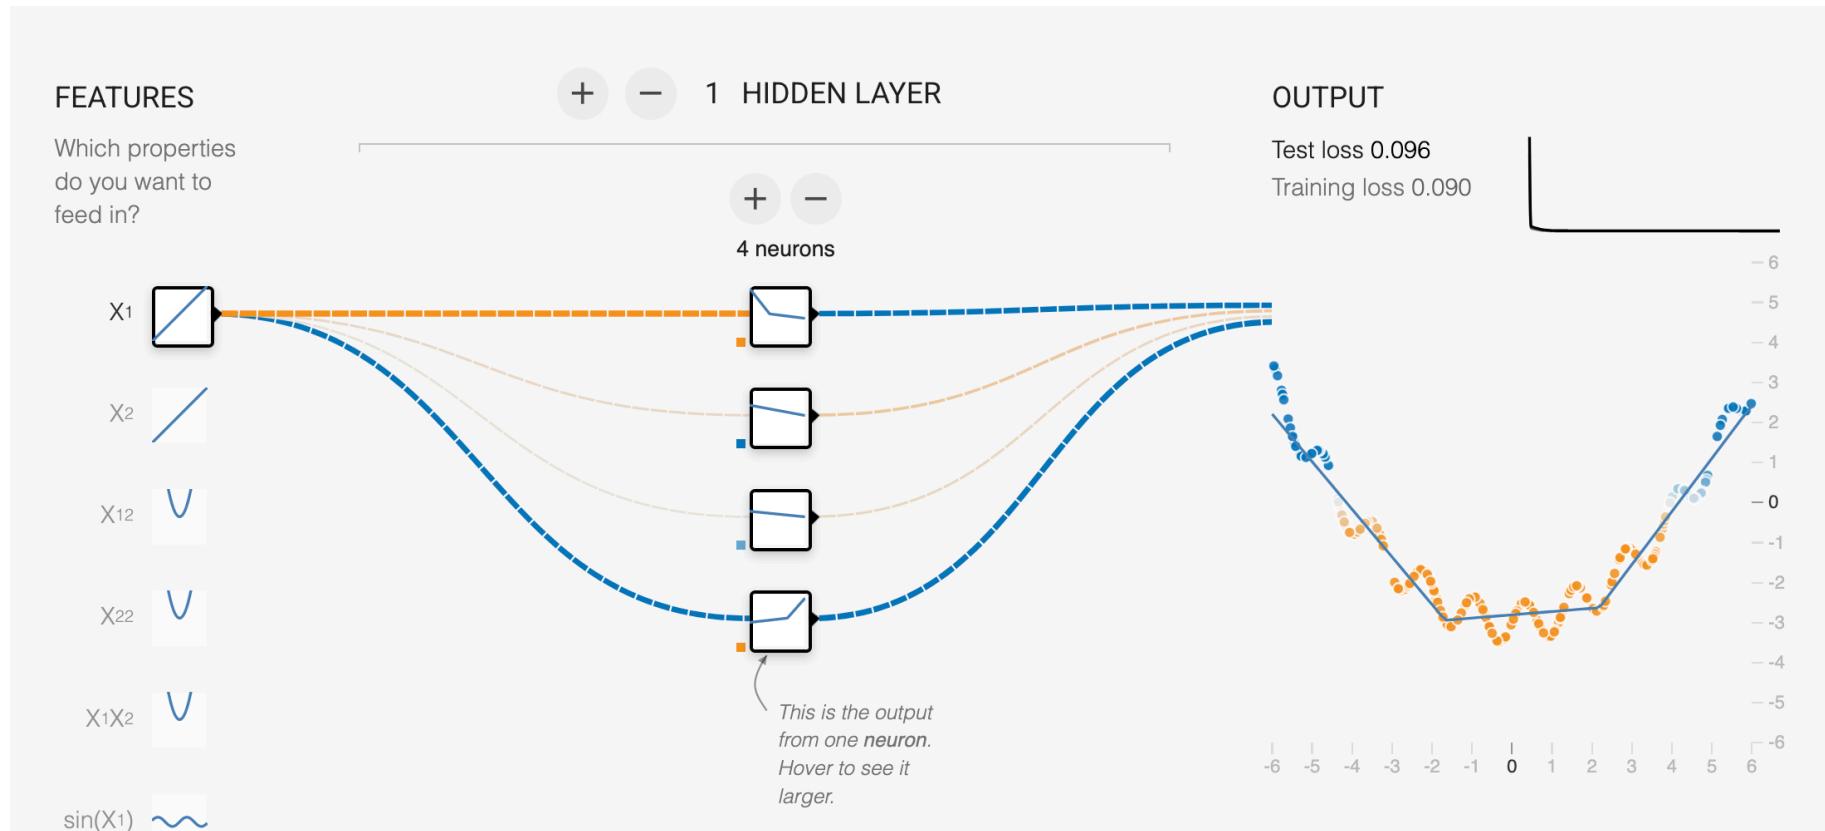
Last time: A Neural Network!



Last time: A Single Neuron



Last time: A Neural Network!

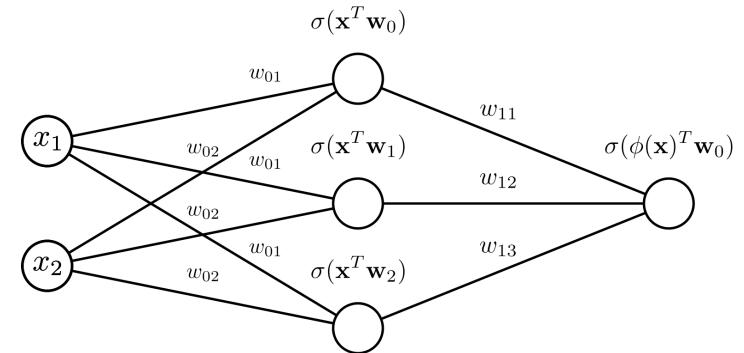


Last time: Neural networks with matrix notation

Our neural network

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}, \quad \mathbf{w}_0 = \begin{bmatrix} w_{01} \\ w_{02} \end{bmatrix}$$

$$f(\mathbf{x}) = \phi(\mathbf{x})^T \mathbf{w}_0, \quad \phi(\mathbf{x}) = \begin{bmatrix} \sigma(\mathbf{x}^T \mathbf{w}_1) \\ \sigma(\mathbf{x}^T \mathbf{w}_2) \\ \sigma(\mathbf{x}^T \mathbf{w}_3) \end{bmatrix} = \begin{bmatrix} \sigma(x_1 w_{11} + x_2 w_{12}) \\ \sigma(x_1 w_{21} + x_2 w_{22}) \\ \sigma(x_1 w_{31} + x_2 w_{32}) \end{bmatrix}$$



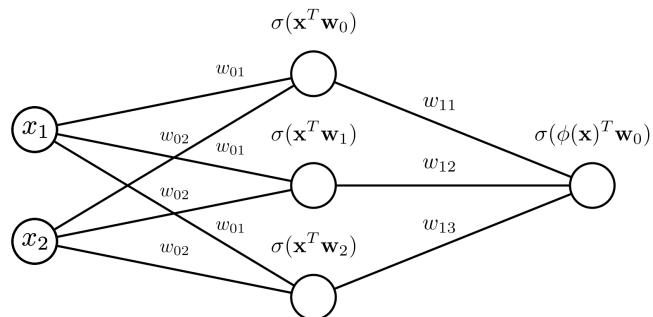
Write the linear function for each neuron as a matrix vector product

$$\begin{bmatrix} \mathbf{x}^T \mathbf{w}_1 \\ \mathbf{x}^T \mathbf{w}_2 \\ \mathbf{x}^T \mathbf{w}_3 \end{bmatrix} = \begin{bmatrix} x_1 w_{11} + x_2 w_{12} \\ x_1 w_{21} + x_2 w_{22} \\ x_1 w_{31} + x_2 w_{32} \end{bmatrix} = \begin{bmatrix} w_{11} & w_{12} \\ w_{21} & w_{22} \\ w_{31} & w_{32} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$$

$$\mathbf{W} = \begin{bmatrix} w_{11} & w_{12} \\ w_{21} & w_{22} \\ w_{31} & w_{32} \end{bmatrix} \quad \begin{bmatrix} \mathbf{x}^T \mathbf{w}_1 \\ \mathbf{x}^T \mathbf{w}_2 \\ \mathbf{x}^T \mathbf{w}_3 \end{bmatrix} = \mathbf{W}\mathbf{x} = (\mathbf{x}^T \mathbf{W}^T)^T$$

Last time: Neural networks with matrix notation

Write all the weights for a layer in a big matrix



$$\mathbf{W}_1 = \begin{bmatrix} \mathbf{w}_1^T \\ \mathbf{w}_2^T \\ \mathbf{w}_3^T \\ \vdots \end{bmatrix} \quad f(\mathbf{x}) = \sigma(\mathbf{W}_1 \mathbf{x})^T \mathbf{w}_0$$

Compact prediction function!

$$\mathbf{X} = \begin{bmatrix} \mathbf{x}_1^T \\ \mathbf{x}_2^T \\ \mathbf{x}_3^T \\ \vdots \end{bmatrix}$$

$$f(\mathbf{x}) = \sigma(\mathbf{X} \mathbf{W}_1) \mathbf{w}_0$$

Can do the same for the full dataset

Neural networks with matrix notation

$$f(\mathbf{x}) = \sigma(\mathbf{X}\mathbf{W}^T)\mathbf{w}_0$$

To summarize:

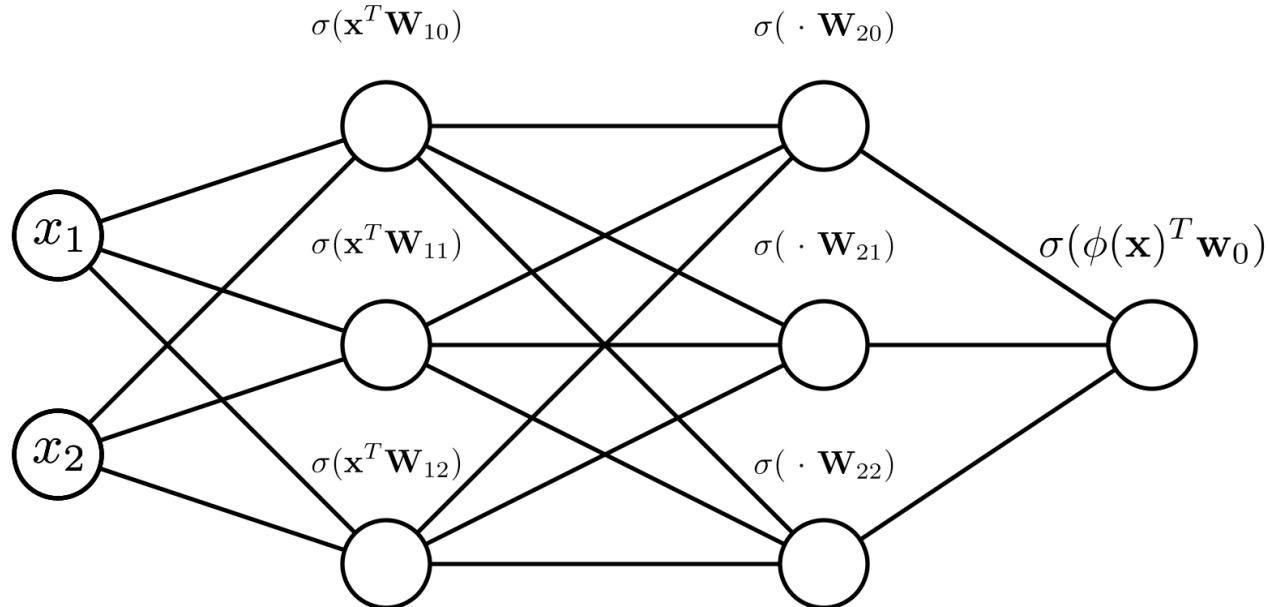
- \mathbf{X} : $N \times d$ matrix of observations
- \mathbf{W} : $h \times d$ matrix of network weights
- \mathbf{w}_0 : h ($\times 1$) vector of linear regression weights

$$\mathbf{X} = \begin{bmatrix} \mathbf{x}_1^T \\ \mathbf{x}_2^T \\ \mathbf{x}_3^T \\ \vdots \end{bmatrix} \quad \mathbf{W}_1 = \begin{bmatrix} \mathbf{w}_1^T \\ \mathbf{w}_2^T \\ \mathbf{w}_3^T \\ \vdots \end{bmatrix} \quad \mathbf{w}_0 = \begin{bmatrix} w_{01} \\ w_{02} \end{bmatrix}$$

If we check that our dimensions work for matrix multiplication we see that we get the $N \times 1$ vector of predictions we are looking for!

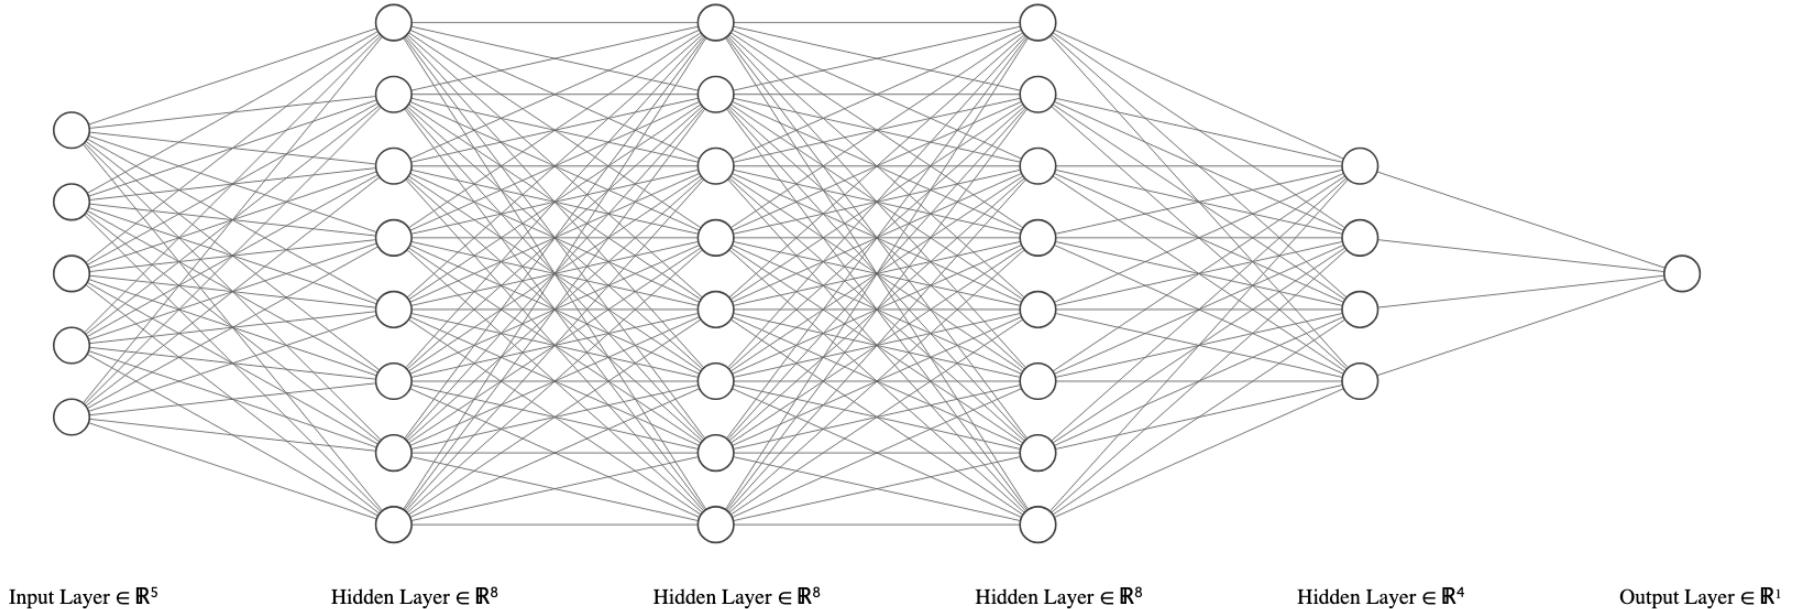
$$(N \times d)(h \times d)^T(h \times 1) \rightarrow (N \times d)(d \times h)(h \times 1) \rightarrow (N \times h)(h \times 1)$$
$$\longrightarrow (N \times 1)$$

Applying this recursively

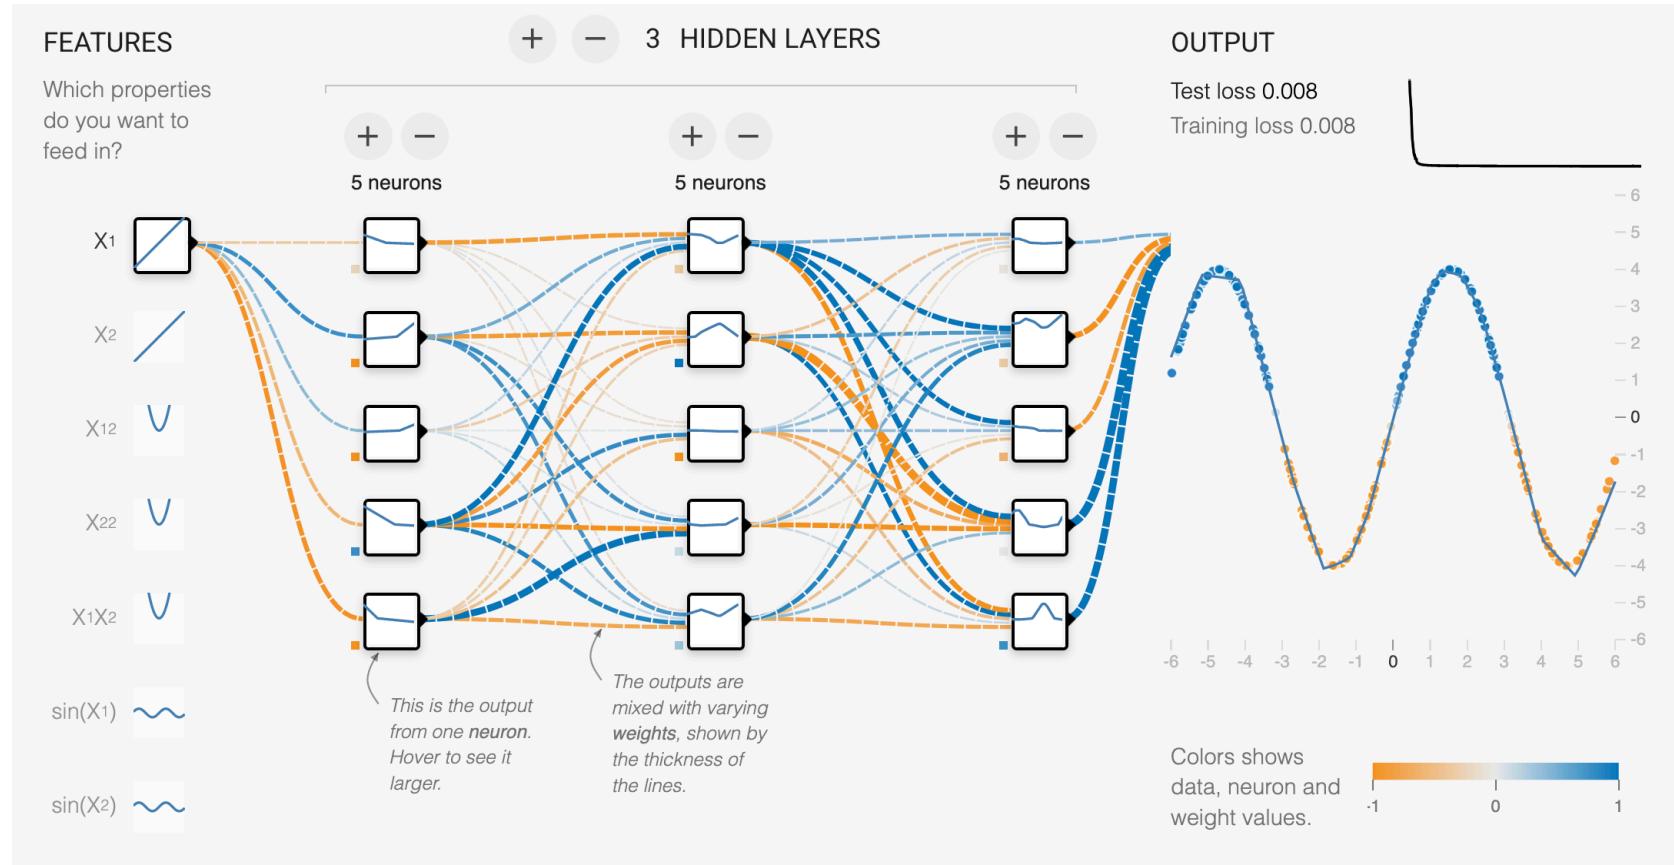


$$\phi(\mathbf{x})_i = \sigma(\mathbf{x}^T \mathbf{w}_i) \quad \phi(\mathbf{x})_i = \sigma(\sigma(\mathbf{x}^T \mathbf{W}^T) \mathbf{w}_i)$$

Applying this recursively



Applying this recursively



Adding Biases

Loss function

$$\begin{aligned}\mathbf{NLL}(\mathbf{w}_0, \mathbf{W}, \mathbf{X}, \mathbf{y}) &= -\sum_{i=1}^N \left[y_i \log p(y = 1 \mid \mathbf{x}, \mathbf{w}_0, \mathbf{W}) + (1 - y_i) \log p(y = 0 \mid \mathbf{x}, \mathbf{w}_0, \mathbf{W}) \right] \\ &= -\sum_{i=1}^N \log \sigma((2y_i - 1)\phi(\mathbf{x}_i)^T \mathbf{w})\end{aligned}$$

Loss function

$$\begin{aligned}\text{NLL}(\mathbf{w}_0, \mathbf{W}, \mathbf{X}, \mathbf{y}) &= -\sum_{i=1}^N \left[y_i \log p(y = 1 \mid \mathbf{x}, \mathbf{w}_0, \mathbf{W}) + (1 - y_i) \log p(y = 0 \mid \mathbf{x}, \mathbf{w}_0, \mathbf{W}) \right] \\ &= -\sum_{i=1}^N \log \sigma((2y_i - 1)\phi(\mathbf{x}_i)^T \mathbf{w})\end{aligned}$$

Substituting in our prediction function

$$p(y = 1 \mid \mathbf{x}, \mathbf{w}_0, \mathbf{W}) = \sigma(\phi(\mathbf{x})^T \mathbf{w}_0) = \sigma(\sigma(\mathbf{x}^T \mathbf{W}^T) \mathbf{w}_0), \quad \phi(\mathbf{x}) = \begin{bmatrix} \sigma(\mathbf{x}^T \mathbf{W}_1) \\ \sigma(\mathbf{x}^T \mathbf{W}_2) \\ \sigma(\mathbf{x}^T \mathbf{W}_3) \end{bmatrix}$$

$$= \sigma(w_{01} \cdot \sigma(x_1 W_{11} + x_2 W_{12}) + w_{02} \cdot \sigma(x_1 W_{21} + x_2 W_{22}) + w_{03} \cdot \sigma(x_1 W_{31} + x_2 W_{32}))$$

Gradient descent

Loss

$$\mathbf{NLL}(\mathbf{w}_0, \mathbf{W}, \mathbf{X}, \mathbf{y}) = - \sum_{i=1}^N \left[y_i \log p(y=1 \mid \mathbf{x}, \mathbf{w}_0, \mathbf{W}) + (1 - y_i) \log p(y=0 \mid \mathbf{x}, \mathbf{w}_0, \mathbf{W}) \right]$$

Updates

$$\mathbf{w}_0^{(k+1)} \leftarrow \mathbf{w}_0^{(k)} - \alpha \nabla_{\mathbf{w}_0} \mathbf{NLL}(\mathbf{w}_0, \mathbf{W}, \mathbf{X}, \mathbf{y}), \quad \mathbf{W}^{(k+1)} \leftarrow \mathbf{W}^{(k)} - \alpha \nabla_{\mathbf{W}} \mathbf{NLL}(\mathbf{w}_0, \mathbf{W}, \mathbf{X}, \mathbf{y})$$

Need:

$$\nabla_{\mathbf{w}_0} \mathbf{NLL}(\mathbf{w}_0, \mathbf{W}, \mathbf{X}, \mathbf{y}) = ? \qquad \qquad \nabla_{\mathbf{W}} \mathbf{NLL}(\mathbf{w}_0, \mathbf{W}, \mathbf{X}, \mathbf{y}) = ?$$

Gradient Shapes

$$\mathbf{NLL}(\mathbf{w}_0, \mathbf{W}, \mathbf{X}, \mathbf{y}) = - \sum_{i=1}^N \left[y_i \log p(y=1 \mid \mathbf{x}, \mathbf{w}_0, \mathbf{W}) + (1 - y_i) \log p(y=0 \mid \mathbf{x}, \mathbf{w}_0, \mathbf{W}) \right]$$

$$\nabla_{\mathbf{w}_0} \mathbf{NLL}(\mathbf{w}_0, \mathbf{W}, \mathbf{X}, \mathbf{y}) = \begin{bmatrix} \frac{\partial \mathbf{NLL}}{\partial w_{01}} \\ \frac{\partial \mathbf{NLL}}{\partial w_{02}} \\ \frac{\partial \mathbf{NLL}}{\partial w_{03}} \\ \vdots \end{bmatrix}, \quad \nabla_{\mathbf{W}} \mathbf{NLL}(\mathbf{w}_0, \mathbf{W}, \mathbf{X}, \mathbf{y}) = \begin{bmatrix} \frac{\partial \mathbf{NLL}}{\partial W_{11}} & \frac{\partial \mathbf{NLL}}{\partial W_{12}} & \cdots & \frac{\partial \mathbf{NLL}}{\partial W_{1d}} \\ \frac{\partial \mathbf{NLL}}{\partial W_{21}} & \frac{\partial \mathbf{NLL}}{\partial W_{22}} & \cdots & \frac{\partial \mathbf{NLL}}{\partial W_{2d}} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial \mathbf{NLL}}{\partial W_{h1}} & \frac{\partial \mathbf{NLL}}{\partial W_{h2}} & \cdots & \frac{\partial \mathbf{NLL}}{\partial W_{hd}} \end{bmatrix}$$

Automatic differentiation

Automatic differentiation review

$$L = -\log \sigma(wx^2)$$

$$\frac{d}{dx} f(g(x)) = f'(g(x))g'(x)$$

$$\begin{aligned} g &= g(x) \\ f &= f(g) \end{aligned} \quad \frac{df}{dx} = \frac{df}{dg} \frac{dg}{dx}$$

3 fundamental steps

1: Break up computation into individual operations

$$a = x^2$$

$$b = wa$$

$$c = \sigma(b)$$

$$g = \log c$$

$$L = -g$$

2: Compute derivative of each step

$$\frac{da}{dx} = 2x \quad \frac{db}{da} = w$$

$$\frac{dc}{db} = \sigma(b)(1 - \sigma(b))$$

$$\frac{dg}{dc} = \frac{1}{c} \quad \frac{dL}{dg} = -1$$

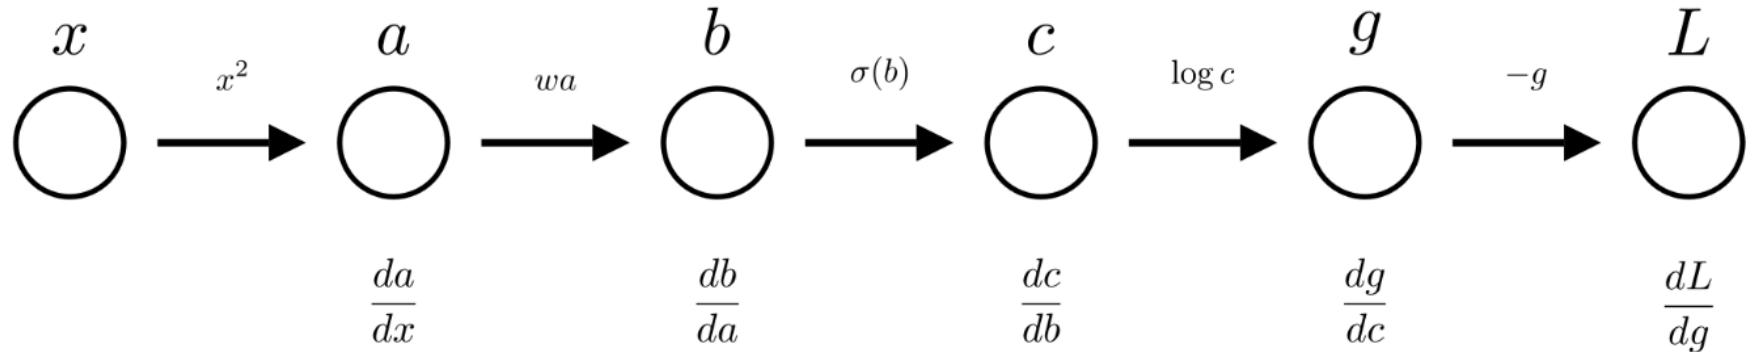
3: Multiply!

$$\frac{dL}{dx} = \frac{dL}{dg} \frac{dg}{dc} \frac{dc}{db} \frac{db}{da} \frac{da}{dx}$$

Computational graph

Nodes with value and derivative for each step

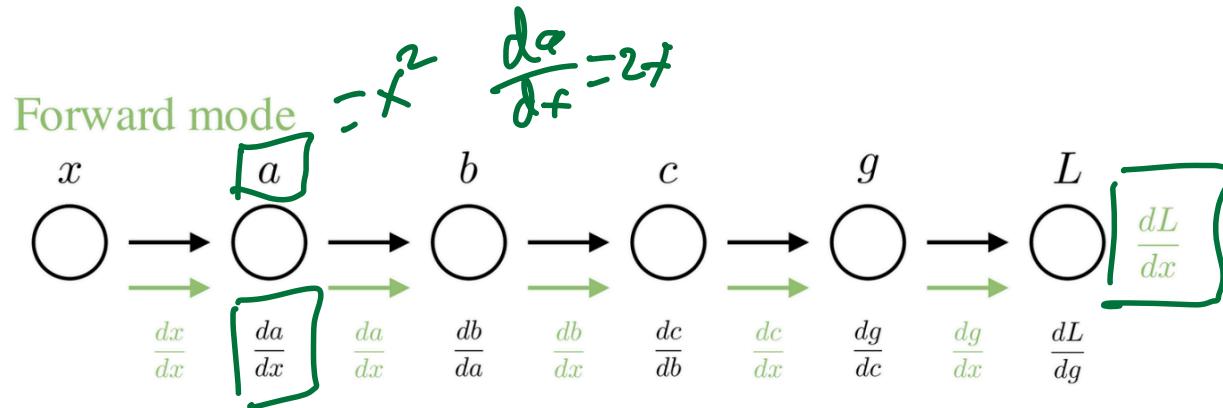
$$L = -\log \sigma(wx^2)$$



Forward-mode autodiff

Keep track of derivative with respect to input
as we go!

$$L = -\log \sigma(wx^2)$$



E.g.: At node **b**

Inputs: a , $\frac{da}{dx}$

Compute:

$b = wa$

$\frac{db}{da} = w$

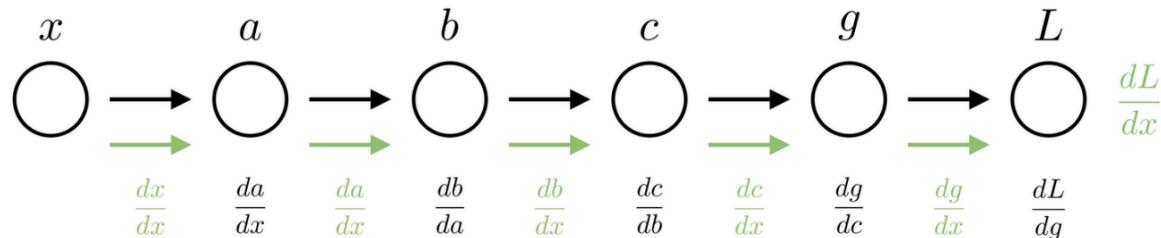
$\frac{d^2b}{dx^2} = \frac{db}{da} \cdot \frac{da}{dx} = w(2x) = 2wx$

Forward-mode autodiff

Keep track of derivative with respect to input
as we go!

$$L = -\log \sigma(wx^2)$$

Forward mode



E.g.: At node b

Inputs: a , $\frac{da}{dx}$

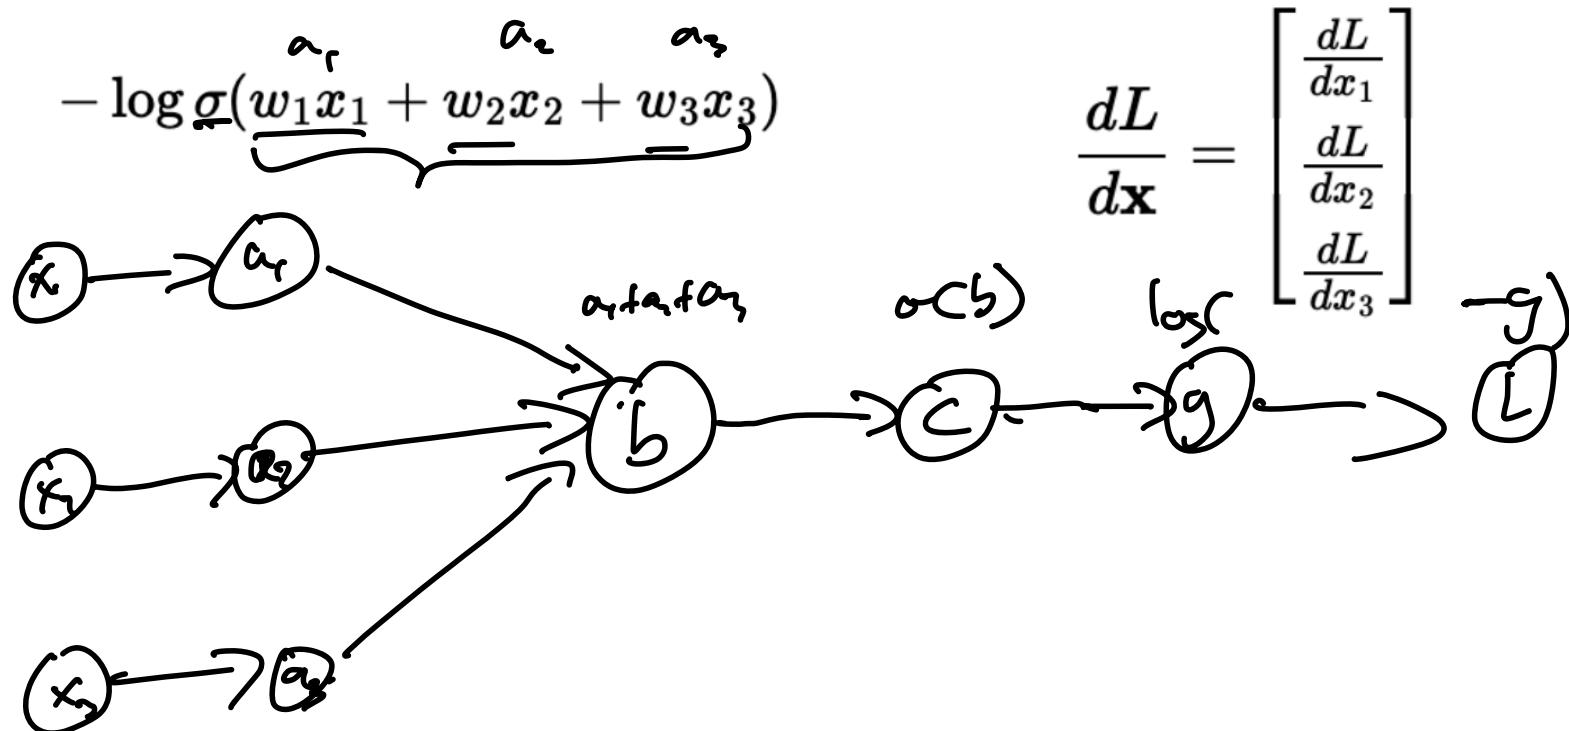
Compute:

$$b = wa \quad \frac{db}{da} = w$$

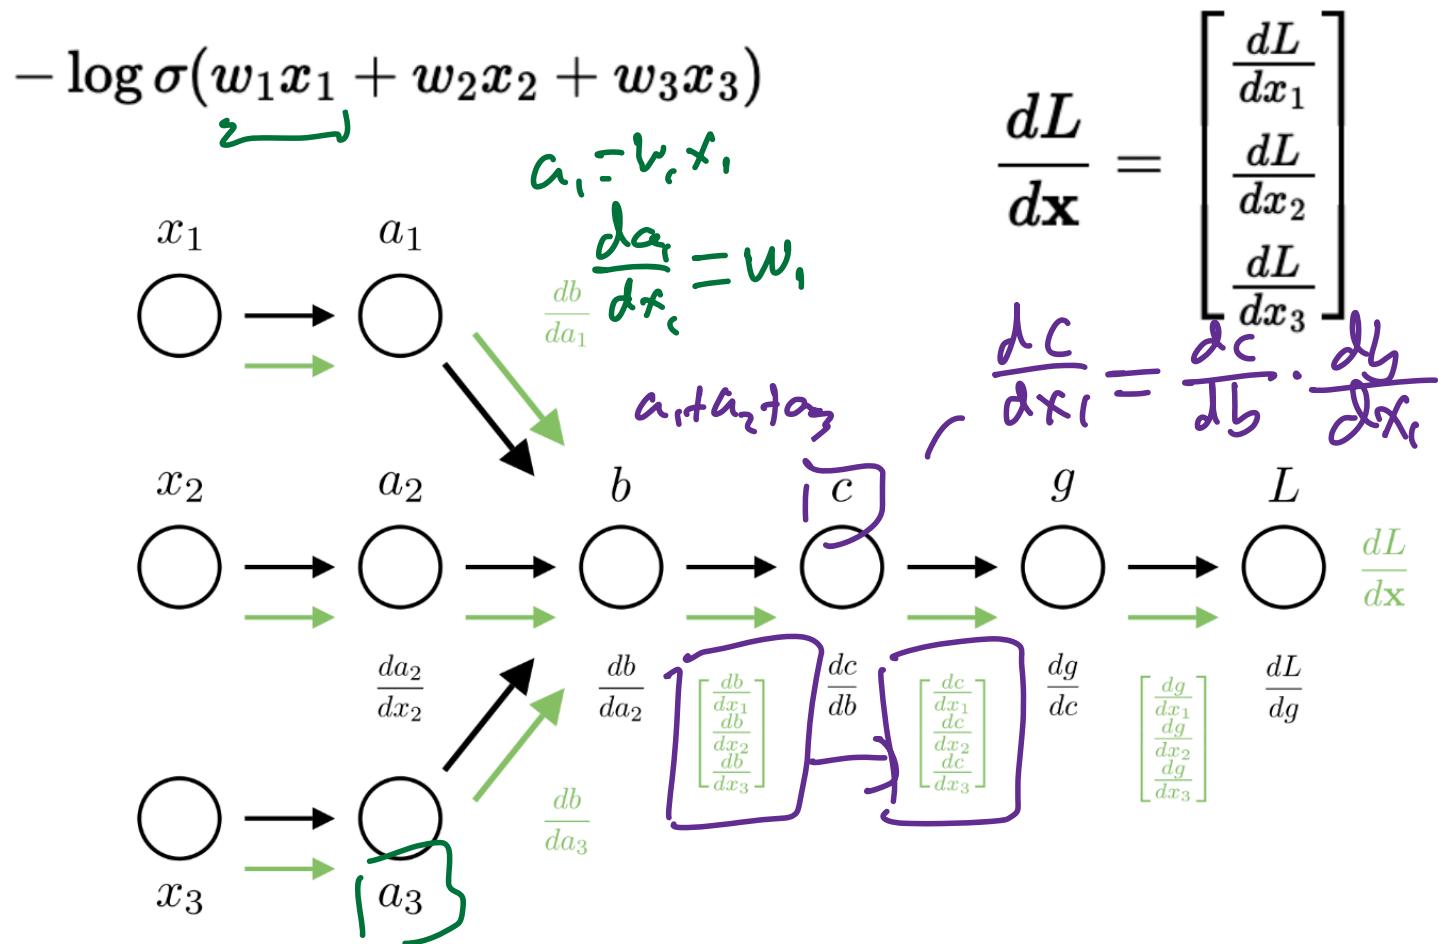
No need to keep!

$$\frac{db}{dx} = \frac{db}{da} \frac{da}{dx}$$

Automatic differentiation with multiple inputs



Automatic differentiation with multiple inputs



Reusing values

```
def loss(x):
    a = x ** 2
    b = 5 * a
    c = log(a)
    g = b * c
    L = -g
    return L
```

$$a = x^2$$

$$b = 5a$$

$$c = \log a$$

$$g = bc$$

$$L = -b$$

Reusing values

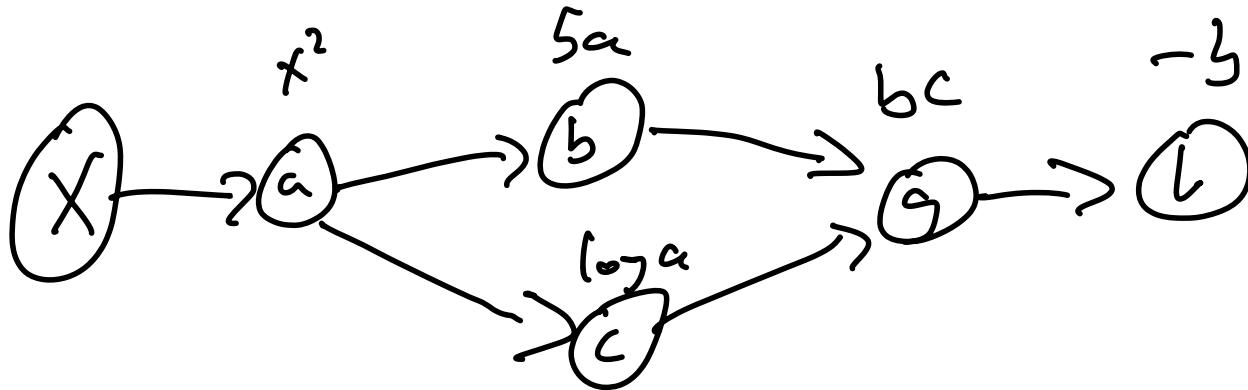
$$a = x^2$$

$$b = 5a$$

$$c = \log a$$

$$g = bc$$

$$L = -\mathfrak{g}$$



b:

$$\frac{dg}{dx} = \frac{dg}{dy} \cdot \frac{dy}{dx}$$

c:

$$\frac{dg}{dx} = \frac{dg}{dc} \cdot \frac{dc}{dx}$$

$$a = x^2$$

$$b = 5a$$

$$c = \log a$$

$$g = bc$$

$$L = -b$$

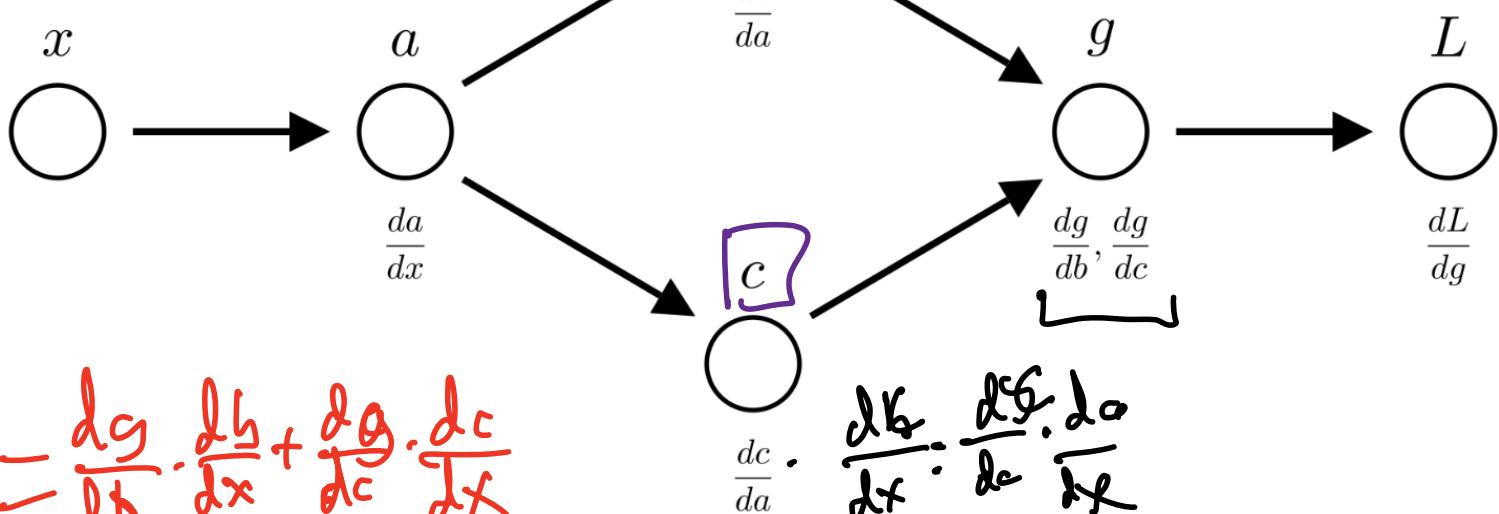
$$\frac{dg}{dx} = \frac{dg}{db} \cdot \frac{db}{dx} + \frac{dg}{dc} \cdot \frac{dc}{dx}$$

Reusing values

$$\frac{dg}{db} = c$$

$$\frac{db}{da} = \frac{db}{dc} \cdot \frac{dc}{da} \quad \frac{dg}{dc} = b$$

$$\frac{dg}{dx} =$$



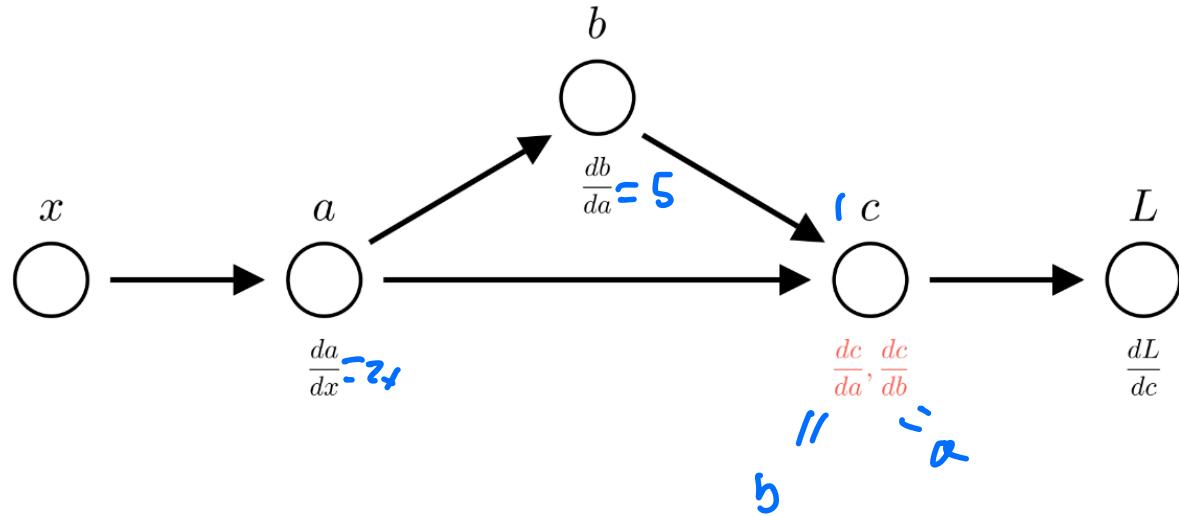
$$\frac{db}{dx} = \frac{dg}{db} \cdot \frac{dg}{dc} \cdot \frac{dc}{da} \cdot \frac{da}{dx}$$

$$L = -\underbrace{(5x^2)}_{x} x^2 = -5x^4$$

Partial derivatives revisited

$$\begin{aligned} a &= x^2 \\ b &= 5a \\ c &= ab \end{aligned}$$

$$L = -c$$



$$\frac{dc}{da} =$$

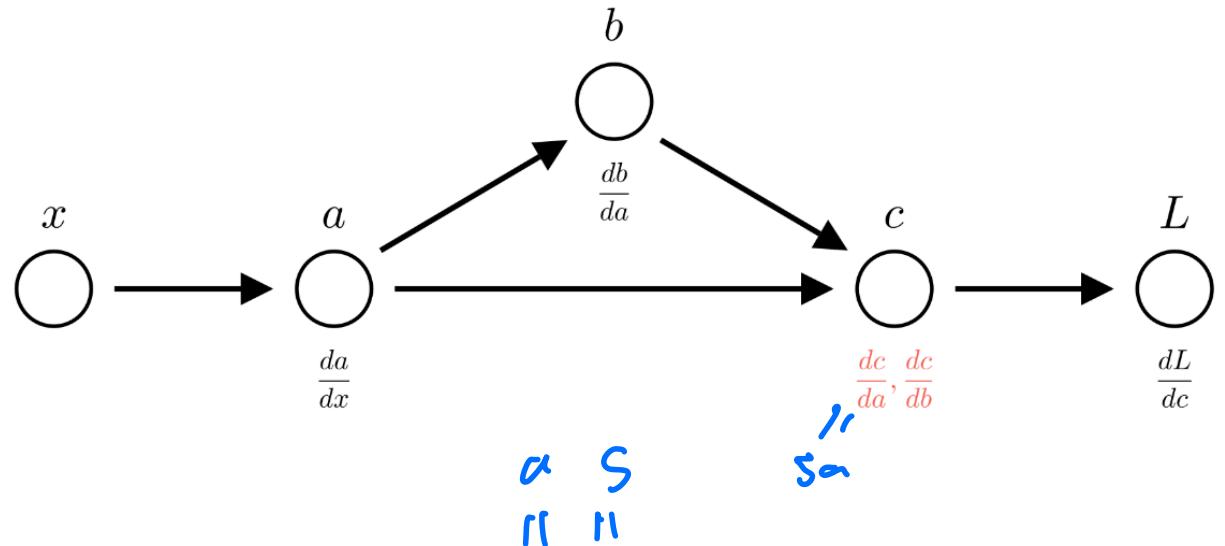
Partial derivatives revisited

$$a = x^2$$

$$b = 5a$$

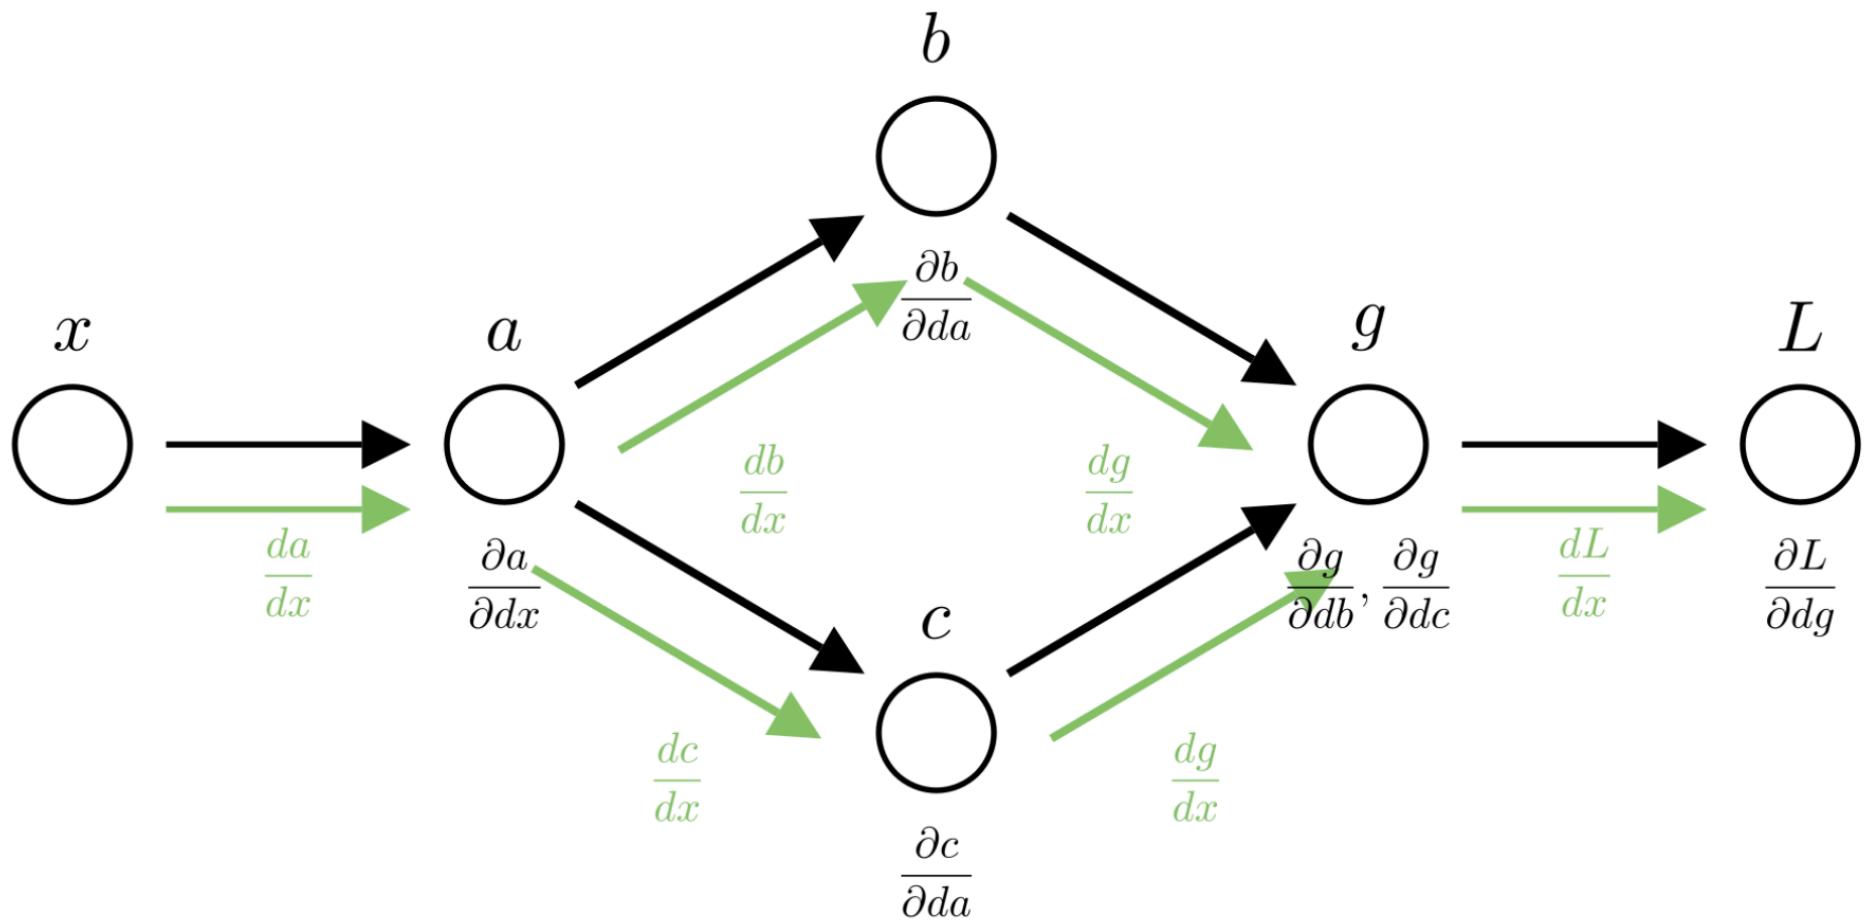
$$c = ab$$

$$L = -c$$

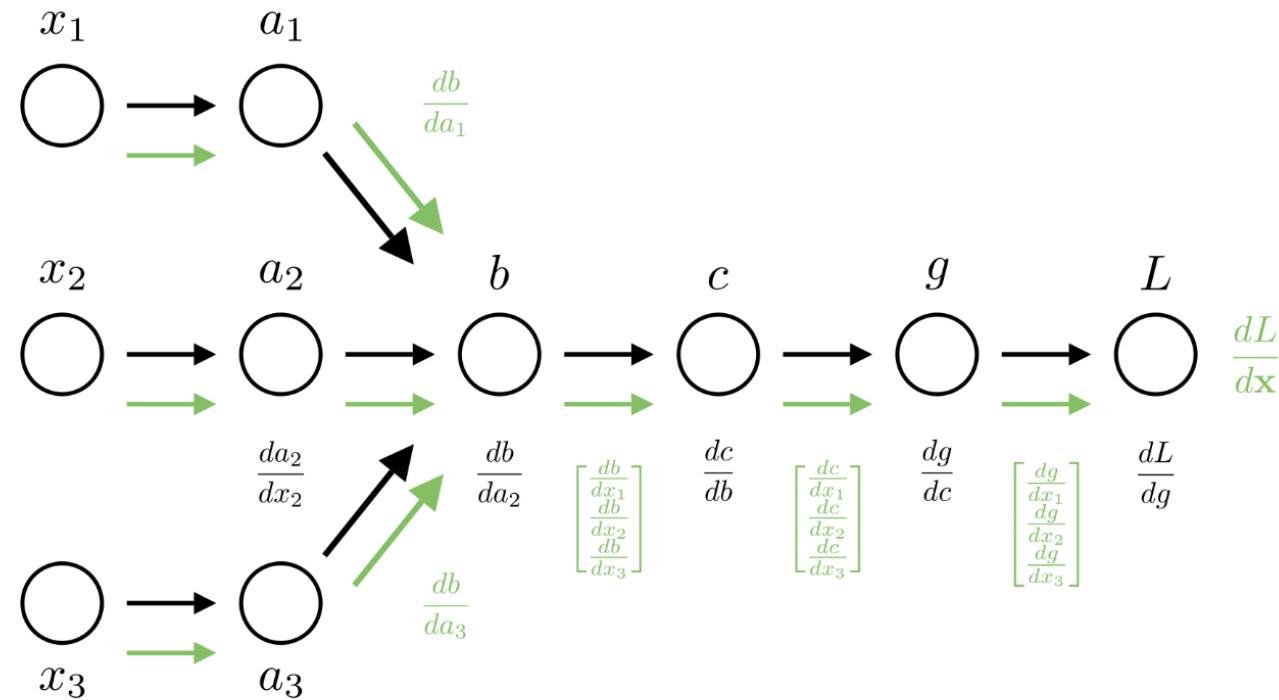


$$\underline{\frac{dc}{da}} = \frac{\partial c}{\partial a} + \frac{\partial c}{\partial b} \underline{\frac{\partial b}{\partial a}} = 5a + 5a = 10a$$

Partial derivatives revisited



Implementing automatic differentiation



```

class ForwardValue:
    ...
    Base class for automatic differentiation operations. Represents variable declaration.
    Subclasses will overwrite func and grads to define new operations.

Properties:
    parent_values (list): A list of raw values of each input (as floats)
    value (float): The value of the result of this operation
    forward_grads (dict): A dictionary mapping inputs to gradients
    ...

def __init__(self, *args):
    self.parent_values = [arg.value if isinstance(arg, ForwardValue) else arg for arg in args]
    self.value = self.forward_pass(args)

    if len(self.forward_grads.keys()) == 0:
        self.forward_grads = {self: 1}

    — how to compute

def func(self, input):
    """
    Compute the value of the operation given the inputs.
    For declaring a variable, this is just the identity function (return the input).

    Args:
        input (float): The input to the operation
    Returns:
        value (float): The result of the operation
    ...
    return input

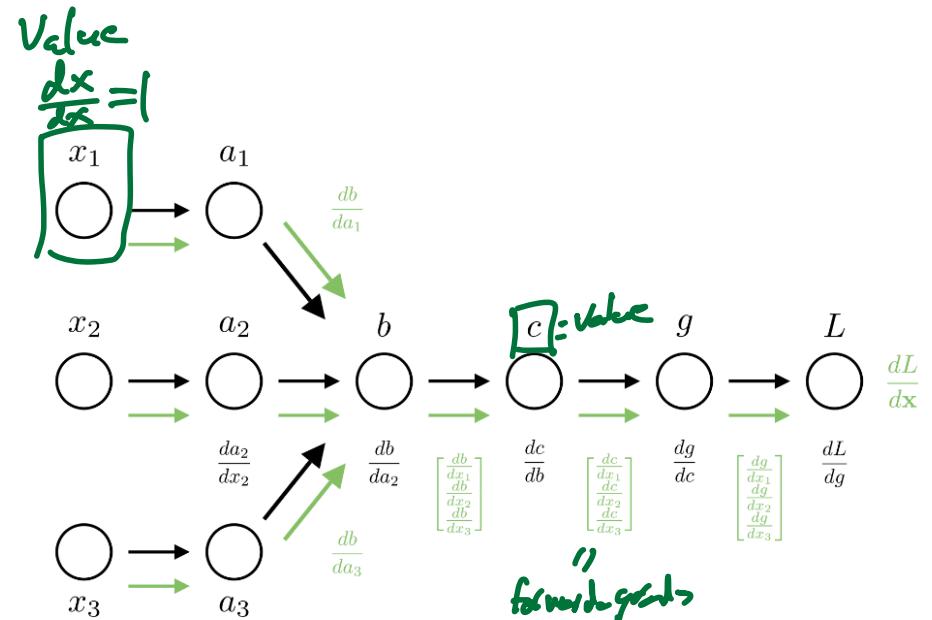
def grads(self, *args):
    """
    Compute the derivative of the operation with respect to each input.
    In the base case the derivative of the identity function is just 1. (da/dx = 1).

    Args:
        input (float): The input to the operation
    Returns:
        grads (tuple): The derivative of the operation with respect to each input
                        Here there is only a single input, so we return a length-1 tuple.
    ...
    return (1,)

def forward_pass(self, args):
    # Calls func to compute the value of this operation
    self.forward_grads = {}
    return self.func(*self.parent_values)

```

Implementing automatic differentiation



$C = a + b$

```
class _add(ForwardValue):
    # Addition operator (a + b)
    def func(self, a, b):
        return a + b

    def grads(self, a, b):
        return 1., 1.
```

$C = a + b$ - result of addition

$$\frac{\partial C}{\partial a} = 1 \quad \frac{\partial C}{\partial b} = 1$$

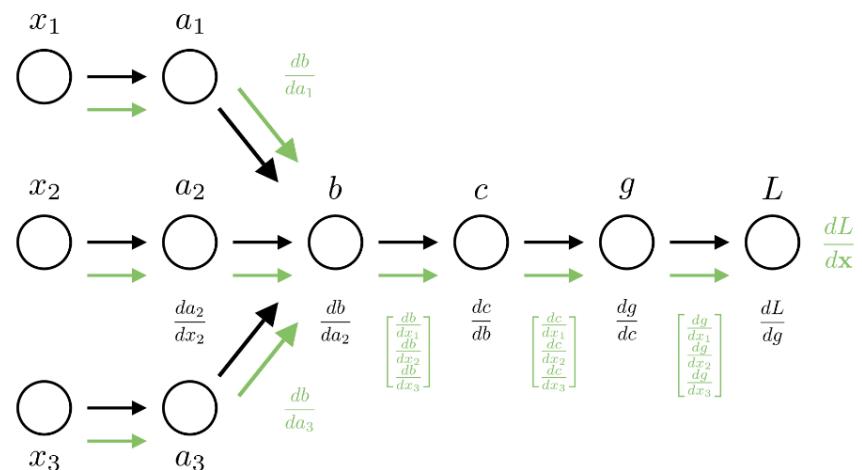
```
class _neg(ForwardValue):
    # Negation operator (-a)
    def func(self, a):
        return -a

    def grads(self, a):
        return (-1.,)

class _sub(ForwardValue):
    # Subtraction operator (a - b)
    def func(self, a, b):
        # Your code here

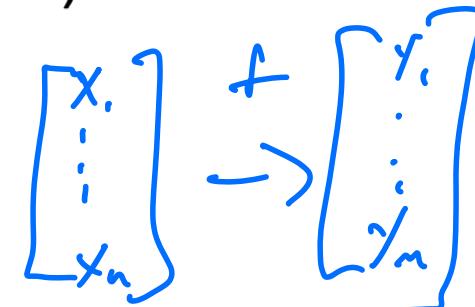
    def grads(self, a, b):
        # Your code here
```

Implementing automatic differentiation



Calculus for vectors (Jacobian)

Now let's consider a vector-valued function $f: \mathbb{R}^n \rightarrow \mathbb{R}^m$
e.g. $f(x_1, x_2, \dots, x_n) = (y_1, y_2, \dots, y_m)$



We define the instantaneous rate of change via the **Jacobian matrix**, whose entries are **partial derivatives**

$$\frac{\partial f}{\partial x} = \begin{bmatrix} \frac{\partial y_1}{\partial x_1} & \frac{\partial y_1}{\partial x_2} & \cdots & \frac{\partial y_1}{\partial x_n} \\ \frac{\partial y_2}{\partial x_1} & \frac{\partial y_2}{\partial x_2} & \cdots & \frac{\partial y_2}{\partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial y_m}{\partial x_1} & \frac{\partial y_m}{\partial x_2} & \cdots & \frac{\partial y_m}{\partial x_n} \end{bmatrix}$$

The Jacobian is also called the **derivative** of f . It is a $m \times n$ matrix

Each partial derivative fixes the other inputs and treats f as scalar-input scalar-output



Calculus for vectors (Jacobian)

m m m / m

$$\mathbf{y} = \mathbf{Ax}$$

$$\frac{dy_j}{dx_i} =$$

$$\frac{\partial \mathbf{y}}{\partial \mathbf{x}} = A \quad \frac{\partial}{\partial x_i} y_j = \sum_{b=1}^n x_b A_{jb} = \frac{1}{\partial x_i} x_i A_{ji} = A_{ji}$$

$$A \begin{bmatrix} \vdots & \vdots \\ \vdots & \vdots \end{bmatrix} [i]_n = []_j$$

$x \in \mathbb{R}^n$

Calculus for vectors (Jacobian)

$$\mathbf{y} = \mathbf{x}^2$$

$$\mathbf{y} = \begin{bmatrix} x_1^2 \\ x_2^2 \\ x_3^2 \\ \vdots \end{bmatrix}$$

$$y_i = x_i^2$$

$$\frac{\partial \mathbf{y}}{\partial \mathbf{x}} = n \begin{bmatrix} 2x_1 & 2x_1 & \square \\ 2x_2 & 2x_2 & \square \\ 2x_3 & 2x_3 & \square \end{bmatrix}$$

$$\frac{\partial y_i}{\partial x_i} = 2x_i$$

$$\frac{\partial y_j}{\partial x_i} = 0$$

 $j \neq i$

Jacobians and gradients

$$f: \mathbb{R}^n \rightarrow \mathbb{R}^m$$

When $m = 1$, (scalar-valued output) the Jacobian is called the **gradient**

$$\frac{\partial f}{\partial x} = \begin{bmatrix} \frac{\partial y}{\partial x_1} & \frac{\partial y}{\partial x_2} & \cdots & \frac{\partial y}{\partial x_n} \end{bmatrix} = \nabla_{\vec{x}} f$$

Notes:

- The Jacobian of a scalar-valued function is a row vector.
- The gradient is often described as a column vector

Chain rule for Jacobians

An easy extension of the chain rule exists:

Given $f: \mathbb{R}^k \rightarrow \mathbb{R}^m$ and $g: \mathbb{R}^n \rightarrow \mathbb{R}^k$

$$y = g(x)$$

$$\frac{\partial f(g(x))}{\partial x} = \underbrace{\frac{\partial f(y)}{\partial y}}_{m \times n} \cdot \underbrace{\frac{\partial g(x)}{\partial x}}_{k \times n}$$

Where $y = g(x)$.

Forward-mode with vectors

$$\mathbf{a} = \mathbf{X}\mathbf{w}$$

$$\mathbf{b} = \mathbf{y} - \mathbf{a}$$

$$\mathbf{c} = \mathbf{b}^2$$

$$g = \sum_{i=1}^N c_i$$

$$L = \frac{1}{N}g$$

Forward-mode with vectors

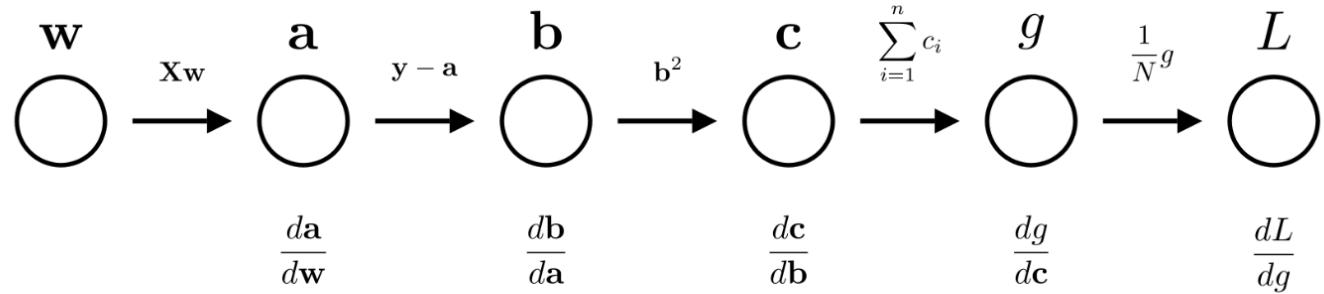
$$\mathbf{a} = \mathbf{X}\mathbf{w}$$

$$\mathbf{b} = \mathbf{y} - \mathbf{a}$$

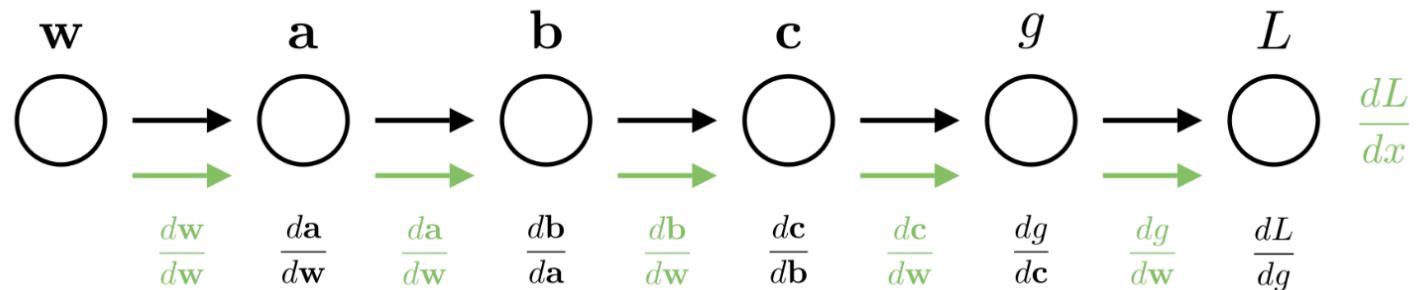
$$\mathbf{c} = \mathbf{b}^2$$

$$g = \sum_{i=1}^N c_i$$

$$L = \frac{1}{N} g$$



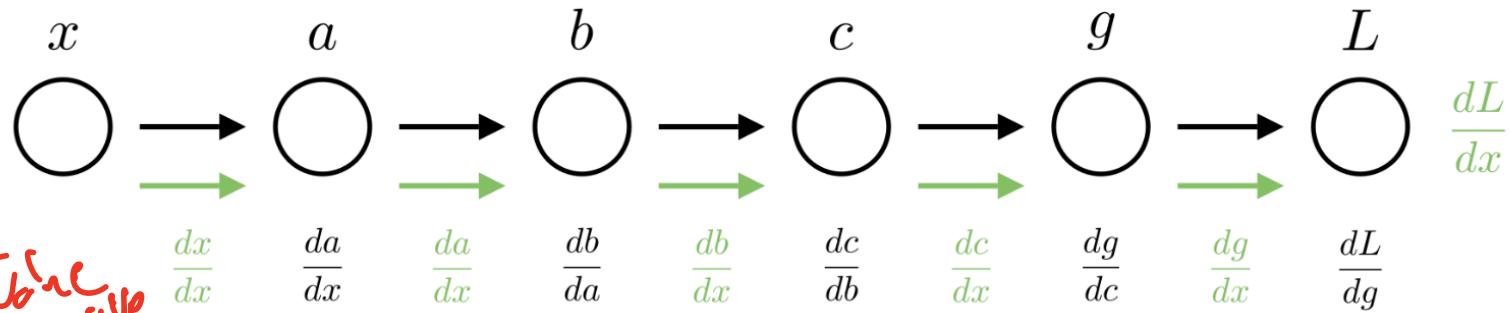
Forward mode



Reverse-mode AD

-pass

Forward mode

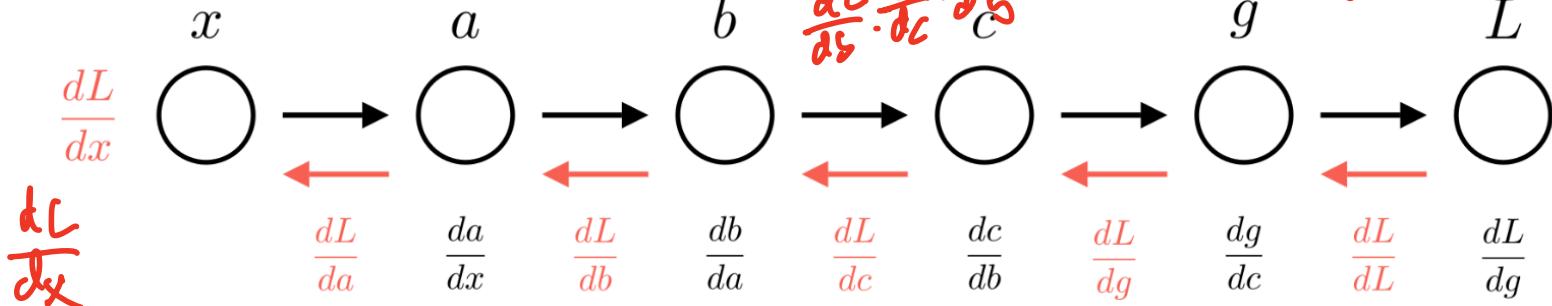


~~1 compact
2 update derivative~~

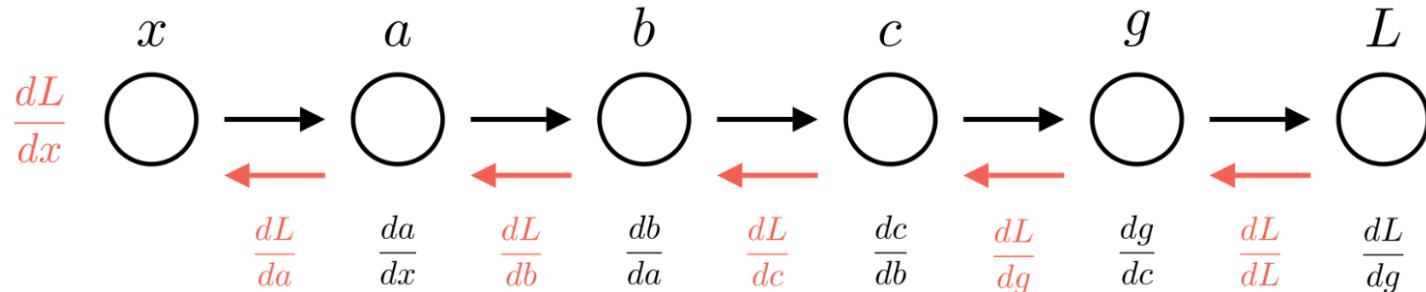
Reverse mode

→ Backpropagation

$$\frac{dL}{dc} = \frac{dL}{dg} \cdot \frac{dg}{dc}$$



Reverse mode



Define inputs

```
a = AutogradValue(5)  
b = AutogradValue(2)
```



Compute operations

```
c = a + b  
L = log(c)
```

Get derivatives

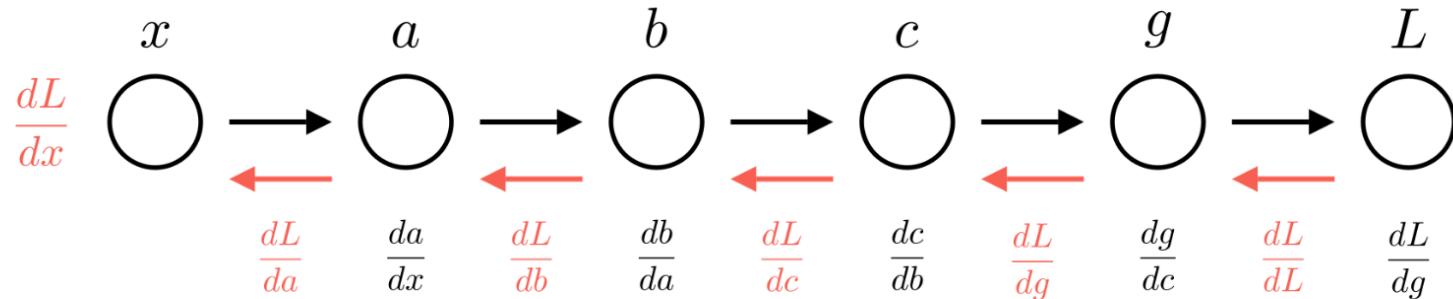
AV

- `L.backward()`
- `dL_da = a.grad`

↑
Stack

Reverse mode

Reverse-mode AD



Define inputs

```
a = AutogradValue(5)  
b = AutogradValue(2)
```

Compute operations

```
c = a + b  
L = log(c)
```

Get derivatives

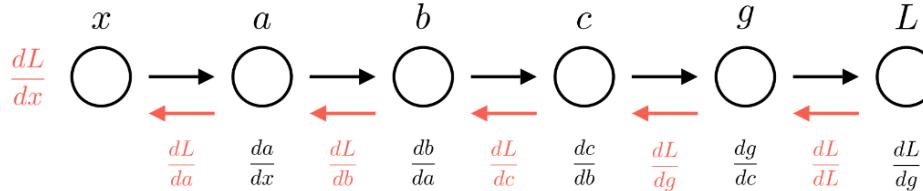
```
L.backward()  
dL_da = a.grad
```

Operations with non-Autograd values

```
s = 4  
L = s * a  
dL_da = a.grad # Will work because a is an AutogradValue  
dL_ds = s.grad # Will give an error because s is not an AutogradValue
```

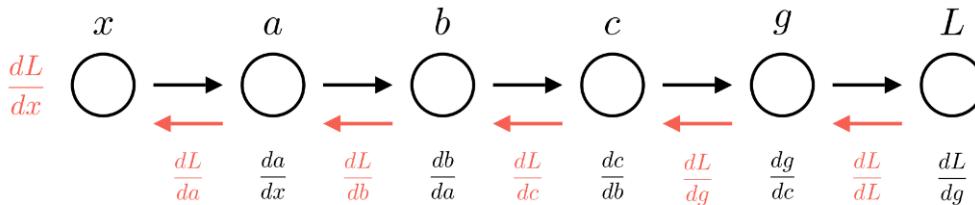
Reverse mode

Reverse-mode AD



```
class AutogradValue:  
    ...  
  
    Base class for automatic differentiation operations.  
    Represents variable delcaration. Subclasses will overwrite  
    func and grads to define new operations.  
  
    Properties:  
        parents (list): A list of the inputs to the operation,  
                        may be AutogradValue or float  
        args      (list): A list of raw values of each  
                        input (as floats)  
        grad      (float): The derivative of the final loss with  
                        respect to this value ( $dL/dx$ )  
        value     (float): The value of the result of this operation  
    ...  
  
    def __init__(self, *args):  
        self.parents = list(args)  
        self.args = [arg.value if isinstance(arg, AutogradValue)  
                    else arg  
                    for arg in self.parents]  
        self.grad = 0.  
        self.value = self.forward_pass()  
  
    def forward_pass(self):  
        # Calls func to compute the value of this operation  
        return self.func(*self.args)
```

Reverse mode



```
class AutogradValue:
    ...
    Base class for automatic differentiation operations.
    Represents variable delcaration. Subclasses will overwrite
    func and grads to define new operations.

Properties:
    parents (list): A list of the inputs to the operation,
                    may be AutogradValue or float
    args    (list): A list of raw values of each
                    input (as floats)
    grad    (float): The derivative of the final loss with
                     respect to this value ( $dL/dx$ )
    value   (float): The value of the result of this operation
    ...

def __init__(self, *args):
    self.parents = list(args)
    self.args = [arg.value if isinstance(arg, AutogradValue)
                else arg
                for arg in self.parents]
    self.grad = 0.
    self.value = self.forward_pass()

def forward_pass(self):
    # Calls func to compute the value of this operation
    return self.func(*self.args)
```

Reverse-mode AD

Base class (input)

```
class AutogradValue:
    def func(self, input):
        ...
        Compute the value of the operation given the inputs.
        For declaring a variable, this is just the identity
        function (return the input).

    Args:
        input (float): The input to the operation
    Returns:
        value (float): The result of the operation
    ...
    return input
```

Sub-class (operation)

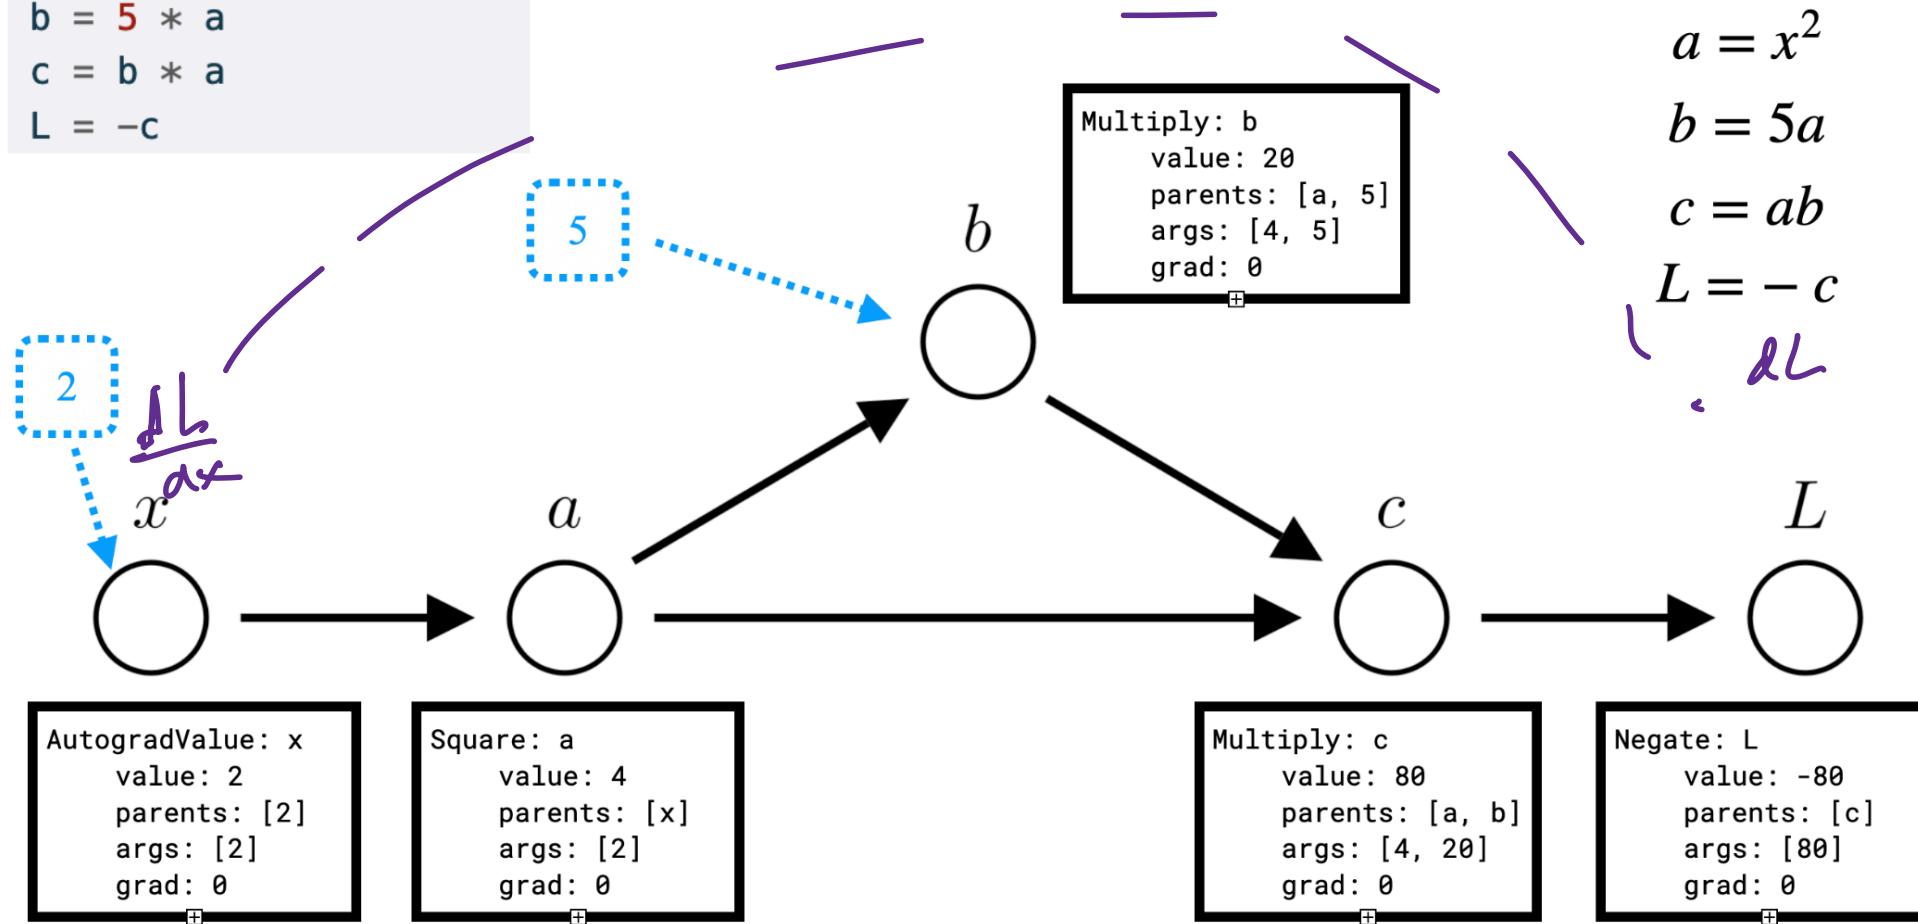
```
class _square(AutogradValue):
    # Square operator ( $a ** 2$ )
    def func(self, a):
        return a ** 2
```

```

x = AutogradValue(2)
a = x ** 2
b = 5 * a
c = b * a
L = -c

```

Forward pass



```
x = AutogradValue(2)
```

```
a = x ** 2
```

```
b = 5 * a
```

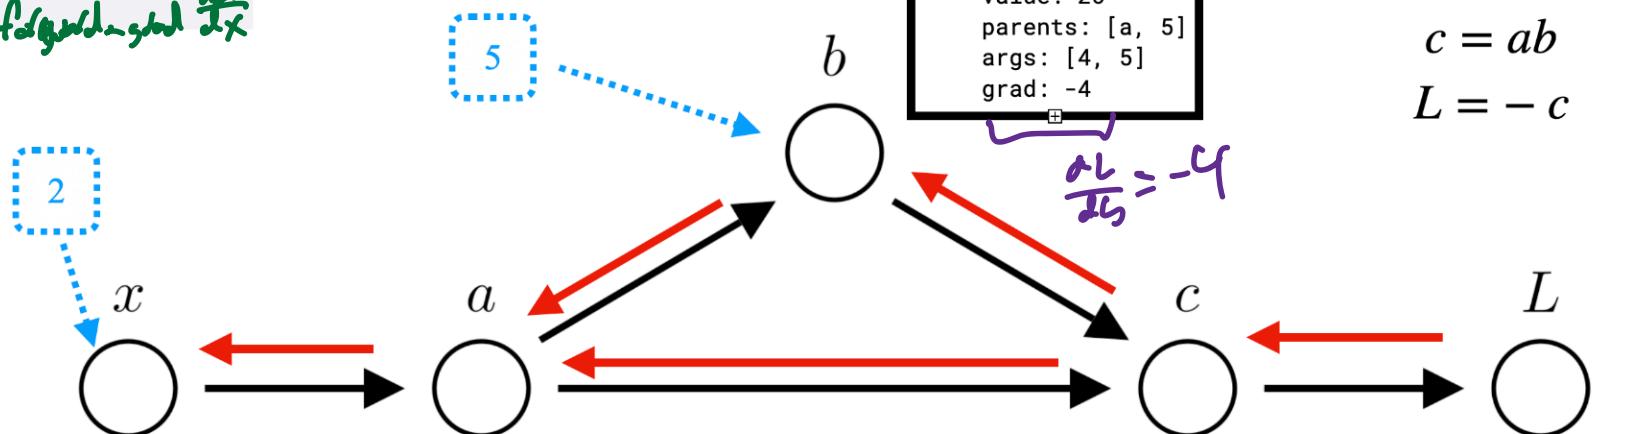
```
c = b * a
```

```
L = -c
```

AutogradValue
value: L
forwarded dx

```
L.backward()  
print('dL_dx', x.grad)
```

Backward pass



AutogradValue: x
value: 2
parents: [2]
args: [2]
grad: -160

Square: a
value: 4
parents: [x]
args: [2]
grad: -40

Multiply: c
value: 80
parents: [a, b]
args: [4, 20]
grad: -1

Negate: L
value: -80
parents: [c]
args: [80]
grad: 1

$$\frac{dL}{dx} = \frac{dL}{da} \frac{da}{dx} = -40 \cdot 4$$

$$\frac{dL}{da} = \frac{dL}{dc} \frac{dc}{da} + \frac{dL}{db} \frac{db}{da} = -1 \cdot 20 + -4 \cdot 5$$

$$\frac{dL}{dc} = -1$$

$$\frac{dL}{dL} = 1$$

5

4

2

1

$$\frac{dL}{dc}$$

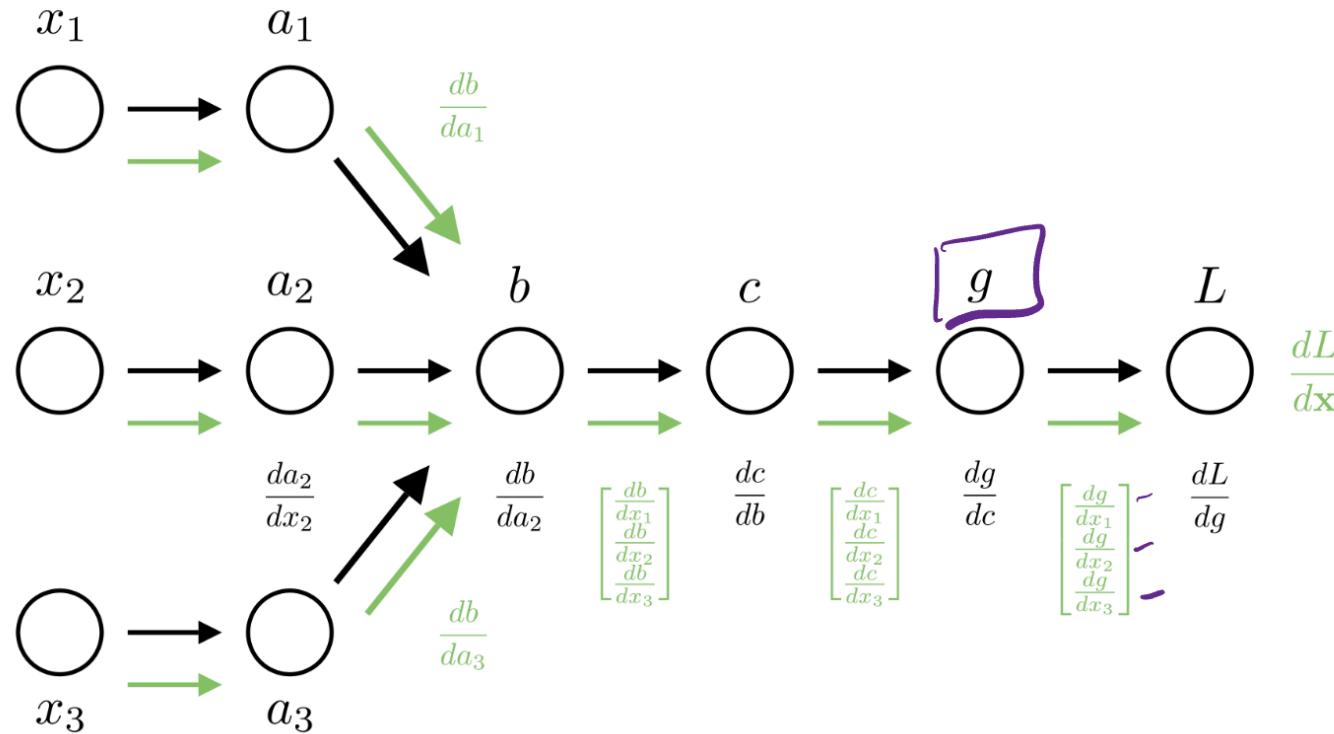
$x = 2$
 $a = x^2$
 $b = 5a$
 $c = ab$
 $L = -c$

$$\frac{dL}{db} = \frac{dL}{dc} \frac{dc}{db} = -1 \cdot 4$$

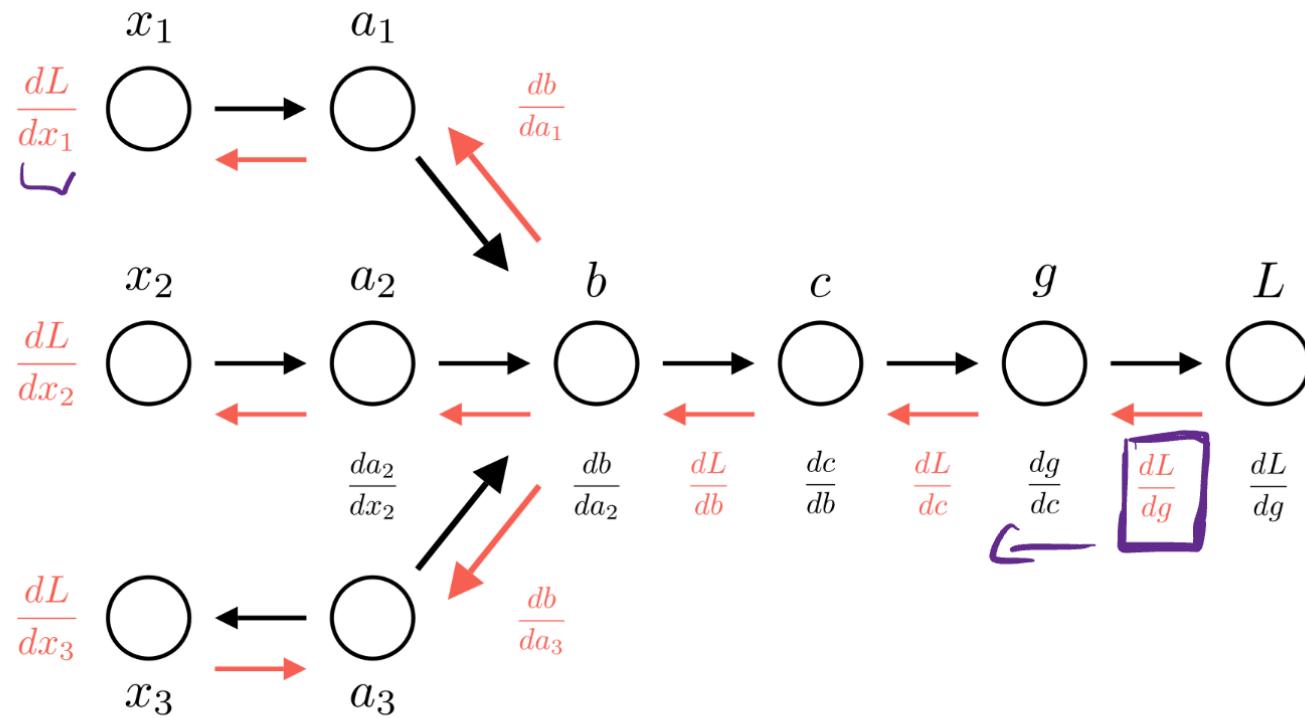
Backward pass

```
class AutogradValue:  
    def grads(self, *args):  
        """  
        Compute the derivative of the operation with respect to each input.  
        In the base case the derivative of the identity function is just 1. (da/da = 1).  
  
        Args:  
            input (float): The input to the operation  
        Returns:  
            grads (tuple): The derivative of the operation with respect to each input  
                           Here there is only a single input, so we return a length-1 tuple.  
        """  
        return (1,)  
  
class _square(AutogradValue):  
    # Square operator (a ** 2)  
    def grads(self, a):  
        return (2 * a,)
```

Why reverse mode?



Why reverse mode?



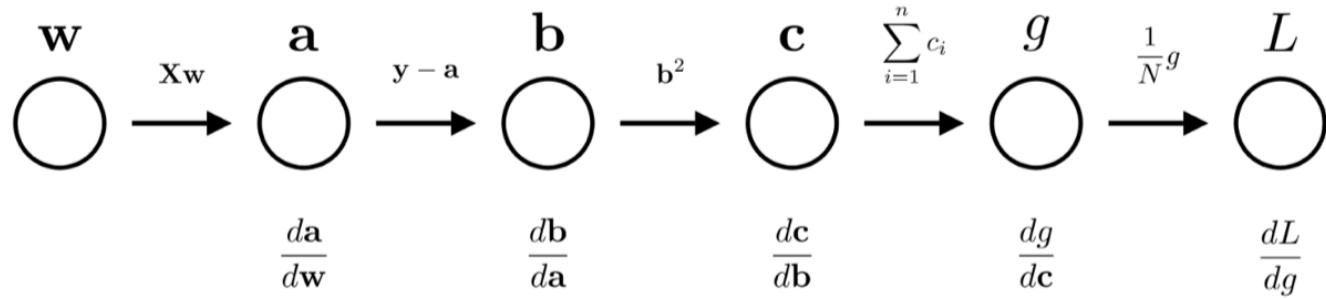
$$\mathbf{a} = \mathbf{Xw}$$

$$\mathbf{b} = \mathbf{y} - \mathbf{a}$$

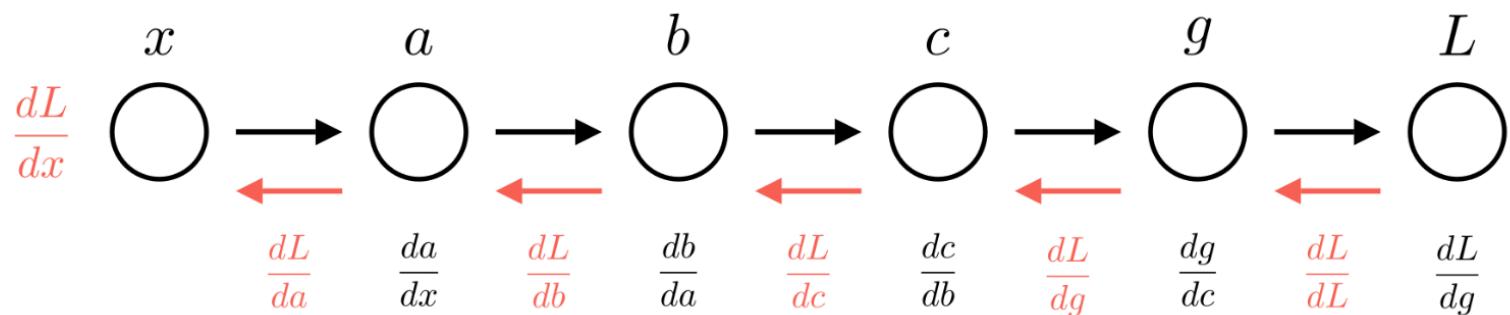
$$\mathbf{c} = \mathbf{b}^2$$

$$g = \sum_{i=1}^N c_i$$

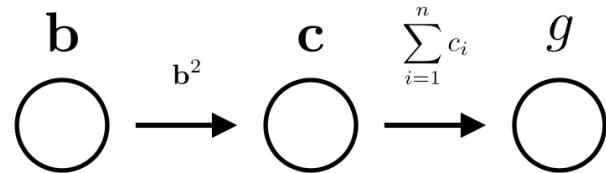
$$L = \frac{1}{N} g$$



Reverse mode



Jacobian computation



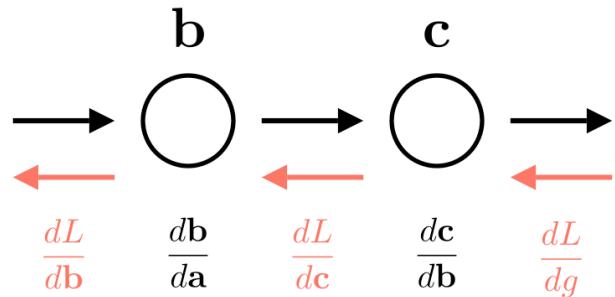
$$\mathbf{c} = \mathbf{b}^2, \quad \mathbb{R}^n \rightarrow \mathbb{R}^n$$

$$\frac{d\mathbf{b}}{da}$$

$$\frac{d\mathbf{c}}{d\mathbf{b}}$$

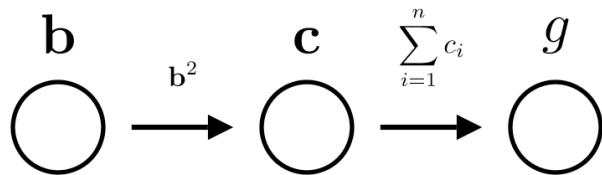
$$\frac{dg}{dc}$$

$$\frac{d\mathbf{c}}{d\mathbf{b}} =$$

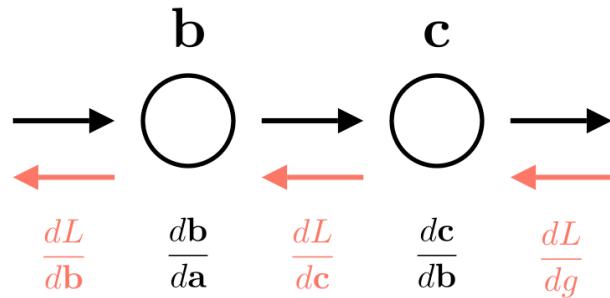


$$\frac{dL}{d\mathbf{b}} = \frac{dL}{dc} \frac{dc}{d\mathbf{b}} =$$

Jacobian computation



$$\frac{db}{da}, \quad \frac{dc}{db}, \quad \frac{dg}{dc}$$

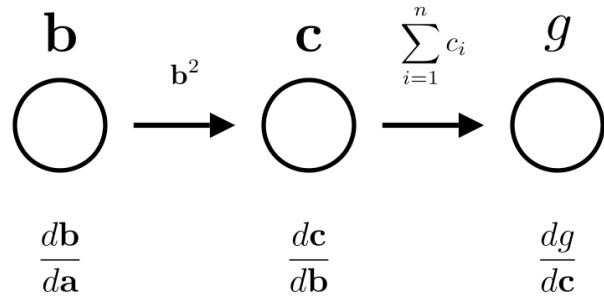


$$\mathbf{c} = \mathbf{b}^2, \quad \mathbb{R}^n \rightarrow \mathbb{R}^n$$

$$\frac{d\mathbf{c}}{d\mathbf{b}} = \begin{bmatrix} \frac{\partial c_1}{\partial b_1} & 0 & \dots & 0 \\ 0 & \frac{\partial c_2}{\partial b_2} & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & \frac{\partial c_N}{\partial b_d} \end{bmatrix}$$

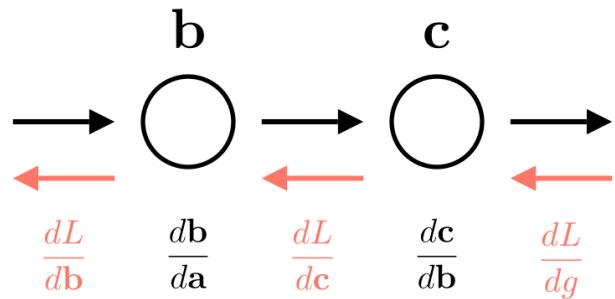
$$\frac{dL}{d\mathbf{b}} = \frac{dL}{d\mathbf{c}} \frac{d\mathbf{c}}{d\mathbf{b}} =$$

Vector-Jacobian Product



$$\mathbf{c} = \mathbf{b}^2, \quad \mathbb{R}^n \rightarrow \mathbb{R}^n$$

$$\frac{d\mathbf{c}}{d\mathbf{b}} = \begin{bmatrix} \frac{\partial c_1}{\partial b_1} & 0 & \dots & 0 \\ 0 & \frac{\partial c_2}{\partial b_2} & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & \frac{\partial c_N}{\partial b_d} \end{bmatrix}$$



$$\frac{dL}{d\mathbf{b}} = \frac{dL}{d\mathbf{c}} \frac{d\mathbf{c}}{d\mathbf{b}} = \begin{bmatrix} \frac{dL}{dc_1} \frac{dc_1}{db_1} \\ \frac{dL}{dc_2} \frac{dc_2}{db_2} \\ \vdots \\ \frac{dL}{dc_N} \frac{dc_N}{db_N} \end{bmatrix}$$

Jacobian computation

$$\mathbf{a} = \mathbf{X}\mathbf{w}, \quad \frac{d\mathbf{a}}{d\mathbf{X}}$$

$$\frac{dL}{d\mathbf{X}} = \frac{dL}{d\mathbf{a}}^T \frac{d\mathbf{a}}{d\mathbf{X}}$$

$$\frac{dL}{dX_{jk}} = \sum_{i=1}^N \frac{\partial L}{\partial a_i} \frac{\partial a_i}{\partial X_{jk}}$$

Jacobian computation

$$\mathbf{a} = \mathbf{X}\mathbf{w}, \quad \frac{d\mathbf{a}}{d\mathbf{X}}$$

$$\frac{dL}{d\mathbf{X}} = \frac{dL}{d\mathbf{a}}^T \frac{d\mathbf{a}}{d\mathbf{X}}$$

```
dL_da, da_dx # The computed gradients/jacobians
da_dx = da_dx.reshape((da_dx.shape[0], -1))
dL_dx = np.dot(dL_da, da_dx)
dL_dx = dL_dx.reshape(x.shape)
```

$$\frac{dL}{dX_{jk}} = \sum_{i=1}^N \frac{\partial L}{\partial a_i} \frac{\partial a_i}{\partial X_{jk}}$$

Jacobian computation

$$\mathbf{a} = \mathbf{X}\mathbf{w}, \quad \frac{d\mathbf{a}}{d\mathbf{X}}$$

$$\frac{dL}{d\mathbf{X}} = \frac{dL}{d\mathbf{a}}^T \frac{d\mathbf{a}}{d\mathbf{X}}$$

$$\frac{dL}{dX_{jk}} = \sum_{i=1}^N \frac{\partial L}{\partial a_i} \frac{\partial a_i}{\partial X_{jk}}$$

$$\mathbf{a} = \mathbf{X}\mathbf{w}, \quad \frac{d\mathbf{a}}{d\mathbf{X}}$$

$$a_i = \sum_{k=1}^d X_{ik} w_k$$

$$\frac{\partial a_i}{\partial X_{jk}} = \frac{\partial}{\partial X_{jk}} \sum_{k=1}^d X_{ik} w_k = \mathbb{I}(i=j) w_k$$

$$\frac{\partial L}{\partial X_{ik}} = \frac{\partial L}{\partial a_i} \frac{\partial a_i}{\partial X_{ik}} = \frac{\partial L}{\partial a_i} w_k$$

Jacobian computation

```
w_mat = w.reshape((1, -1)) # reshape w to 1 x d  
dL_da_mat = dL_da.reshape((-1, 1)) # reshape dL_da to N x 1  
dL_dx = w_mat * dL_da_mat # dL_dx becomes N x d, entry ik = dL_da_i * w_k
```

```
dL_da, da_dx # The computed gradients/jacobians  
da_dx = da_dx.reshape((da_dx.shape[0], -1))  
dL_dx = np.dot(dL_da, da_dx)  
dL_dx = dL_dx.reshape(x.shape)
```

$$\mathbf{a} = \mathbf{X}\mathbf{w}, \quad \frac{d\mathbf{a}}{d\mathbf{X}}$$

$$a_i = \sum_{k=1}^d X_{ik} w_k$$

$$\frac{\partial a_i}{\partial X_{jk}} = \frac{\partial}{\partial X_{jk}} \sum_{k=1}^d X_{ik} w_k = \mathbb{I}(i=j) w_k$$

$$\frac{\partial L}{\partial X_{ik}} = \frac{\partial L}{\partial a_i} \frac{\partial a_i}{\partial X_{ik}} = \frac{\partial L}{\partial a_i} w_k$$

Vector-Jacobian products

$$f(\mathbf{x}) = \mathbf{x} + \mathbf{a}$$

$$\mathbf{v}^T \frac{d\mathbf{f}}{d\mathbf{x}} =$$

$$f(\mathbf{x}) = \log \mathbf{x}$$

$$\mathbf{v}^T \frac{d\mathbf{f}}{d\mathbf{x}} =$$